Cho, J., Garcia-Molina, H., and Page, L. (1998). Efficient crawling through URL ordering. *Comp. Netw.*, 30(161–172). Proceedings of WWW Conference. 664, 665

Cho, J. and Ntoulas, A. (2002). Effective change detection using sampling. In *Proc. 28th Int. Conf. on Very Large Data Bases*. 666

Chockler, G., Keidar, I., and Vitenberg, R. (2001). Group communication specifications: a comprehensive study. *ACM Comput. Surv.*, 33(4):427–469. 482, 537

Christensen, E., Curbera, F., Meredith, G., and Weerawarana, S., editors. Web services description language (WSDL) 1.1 (2001). Available from: `http://www.w3.org/TR/wsdl` [Last retrieved: December 2009]. 690

Chu, W. W. (1969). Optimal file allocation in a multiple computer system. *IEEE Trans. Comput.*, C-18(10):885–889. 125

Chu, W. W. (1973). Optimal file allocation in a computer network. In Abramson, N. and Kuo, F. F., editors, *Computer Communication Networks*, pages 82–94. Prentice-Hall. 125

Chu, W. W. (1976). Performance of file directory systems for data bases in star and distributed networks. In *Proc. National Computer Conf*, pages 577–587. 38

Chu, W. W. and Nahouraii, E. E. (1975). File directory design considerations for distributed databases. In *Proc. 1st Int. Conf. on Very Data Bases*, pages 543–545. 38

Chundi, P., Rosenkrantz, D. J., and Ravi, S. S. (1996). Deferred updates and data placement in distributed databases. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 469–476. 477

Civelek, F. N., Dogac, A., and Spaccapietra, S. (1988). An expert system approach to view definition and integration. In *Proc. 7th Int'l. Conf. on Entity-Relationship Approach*, pages 229–249. 202

Clarke, I., Miller, S. G., Hong, T. W., Sandberg, O., and Wiley, B. (2002). Protecting free expression online with Freenet. *IEEE Internet Comput.*, 6(1):40–49. 615

Clarke, I., Sandberg, O., Wiley, B., and Hong, T. W. (2000). Freenet: A distributed anonymous information storage and retrieval system. In *Proc. Workshop on Design Issues in Anonymity and Unobservability*, pages 46–66. 615

Cluet, S. and Delobel, C. (1992). A general framework for the optimization of object-oriented queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 383–392. 583, 586, 587, 588

Codd, E. (1995). Twelve rules for on-line analytical processing. *Computerworld*. 132

Codd, E. F. (1970). A relational model for large shared data banks. *Commun. ACM*, 13(6):377–387. 45, 56

Codd, E. F. (1972). Relational completeness of data base sublanguages. In Rustin, R., editor, *Relational Databases*, pages 65–98. Prentice-Hall, Englewood Cliffs, N.J. 45

Codd, E. F. (1974). Recent investigations in relational data base systems. *Proceedings of IFIP Congress, Information Processing 74*, pages 1017–1021. 44

Codd, E. F. (1979). Extending the database relational model to capture more meaning. *ACM Trans. Database Syst.*, 4(4):397–434. 43

Cohen, E. and Kaplan, H. (2004). Spatially-decaying aggregation over a network: Model and algorithms. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 707–718. 726

Cohen, E. and Strauss, M. (2003). Maintaining time-decaying stream aggregates. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 223–233. 726, 737

Cohen, S. (2006). User-defined aggregate functions: bridging theory and practice. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 49–60. 737

Cole, R. L. and Graefe, G. (1994). Optimization of dynamic query evaluation plans. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 150–160. 265, 266, 292

Colouris, G., Dollimore, J., and Kindberg, T. (2001). *Distributed Systems: Concepts and Design*. Addison Wesley, 3 edition. 2

Comer, D. E. (2009). *Computer Networks and Internets*. Prentice-Hall, 5 edition. 70

Computers, S. (1982). *Stratus/32 System Overview*. Stratus, Natick, Mass. 456

Cong, G., Fan, W., and Kementsietsidis, A. (2007). Distributed query evaluation with performance guarantees. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 509–520. 711

Cooper, B. F., Ramakrishnan, R., Srivastava, U., Silberstein, A., Bohannon, P., Jacobsen, H.-A., Puz, N., Weaver, D., and Yerneni, R. (2008). PNUTS: Yahoo!'s hosted data serving platform. *Proc. VLDB*, 1(2):1277–1288. 757, 763

Copeland, G., Alexander, W., Boughter, E., and Keller, T. (1988). Data placement in bubba. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 99–108. 510, 511

Copeland, G. and Maier, D. (1984). Making smalltalk a database system. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 316–325. 552

Cormode, G. and Muthukrishnan, S. (2003). What's hot and what's not: Tracking most frequent items dynamically. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 296–306. 743

Coulon, C., Pacitti, E., and Valduriez, P. (2005). Consistency management for partial replication in a high performance database cluster. In *Proc. IEEE Int. Conf. on Parallel and Distributed Systems*, pages 809–815. 537, 539, 548

Crainiceanu, A., Linga, P., Gehrke, J., and Shanmugasundaram, J. (2004). Querying peer-to-peer networks using p-trees. In *Proc. 7th Int. Workshop on the World Wide Web and Databases*, pages 25–30. 622

Cranor, C., Johnson, T., Spatscheck, O., and Shkapenyuk, V. (2003). Gigascope: High performance network monitoring with an SQL interface. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 647–651. 728, 731

Crespo, A. and Garcia-Molina, H. (2002). Routing indices for peer-to-peer systems. In *Proc. 22nd Int. Conf. on Distributed Computing Systems*, pages 23–33. 617

Cristian, F. (1982). Exception handling and software fault tolerance. *IEEE Trans. Comput.*, C-31(6):531–540. 455

Cristian, F. (1985). A rigorous approach to fault–tolerant programming. *IEEE Trans. Softw. Eng.*, SE-11(1):23–31. 455

Cristian, F. (1987). Exception handling. Technical Report RJ 5724, IBM Almaden Research Laboratory, San Jose, Calif. 455

Cuenca-Acuna, F., Peery, C., Martin, R., and Nguyen, T. (2003). Planetp: using gossiping to build content addressable peer-to-peer information sharing communities. In *IEEE Int. Symp. on High Performance Distributed Computing*, pages 236–249. 636

Cusumano, M. A. (2010). Cloud computing and SaaS as new computing platforms. *Commun. ACM*, 53(4):27–29. 744, 763

Dadam, P. and Schlageter, G. (1980). Recovery in distributed databases based on non-synchronized local checkpoints. In *Information Processing '80*, pages 457–462. 456

Dageville, B., Casadessus, P., and Borla-Salamet, P. (1994). The impact of the ksr1 allcache architecture on the behavior of the dbs3 parallel dbms. In *Proc. International Conf. on Parallel Architectures and Language*. 528, 548

Dahlin, M., Wang, R., Anderson, T., and Patterson, D. (1994). Cooperative caching: Using remote client memory to improve file system performance. In *Proc. 1st USENIX Symp. on Operating System Design and Implementation*, pages 267–280. 210

Das, A., Gehrke, J., and Riedewald, M. (2005). Semantic approximation of data stream joins. *IEEE Trans. Knowl. and Data Eng.*, 17(1):44–59. 740

Dasu, T., Krishnan, S., Venkatasubramanian, S., and Yi, K. (2006). An information-theoretic approach to detecting changes in multi-dimensional data streams. In *Proc. 38th Symp. on the Interface of Stats, Comp. Sci., and Applications*. 727

Daswani, N., Garcia-Molina, H., and Yang, B. (2003). Open problems in data-sharing peer-to-peer systems. In *Proc. 9th Int. Conf. on Database Theory*, pages 1–15. 611, 653

Datar, M., Gionis, A., Indyk, P., and Motwani, R. (2002). Maintaining stream statistics over sliding windows. In *Proc. 13th Annual ACM-SIAM Symp. on Discrete Algorithms*, pages 635–644. 737

Date, C. and Darwen, H. (1998). *Foundation for Object/Relational Databases – The Third Manifesto*. Addison Wesley. 552, 607

Date, C. J. (1987). *A Guide to the SQL Standard*. Addison Wesley. 56

Date, C. J. (2004). *An Introduction to Database Systems*. Pearson, 8th edition. 70

Daudjee, K. and Salem, K. (2004). Lazy database replication with ordering guarantees. In *Proc. 20th Int. Conf. on Data Engineering*, pages 424–435. 466

Daudjee, K. and Salem, K. (2006). Lazy database replication with snapshot isolation. In *Proc. 32nd Int. Conf. on Very Large Data Bases*, pages 715–726. 464, 466

Davenport, R. A. (1981). Design of distributed data base systems. *Comp. J.*, 24(1):31–41. 73

Davidson, S. B. (1984). Optimism and consistency in partitioned distributed database systems. *ACM Trans. Database Syst.*, 9(3):456–481. 456, 487, 493

Davidson, S. B., Garcia-Molina, H., and Skeen, D. (1985). Consistency in partitioned networks. *ACM Comput. Surv.*, 17(3):341–370. 449, 456, 493

Dawson, J. L. (1980). A user demand model for distributed database design. In *Digest of Papers – COMPCON*, pages 211–216. 125

Dayal, U. (1989). Queries and views in an object-oriented data model. In *Proc. 2nd Int. Workshop on Database Programming Languages*, pages 80–102. 555, 606

Dayal, U. and Bernstein, P. (1978). On the updatability of relational views. In *Proc. 4th Int. Conf. on Very Data Bases*, pages 368–377. 175, 201

Dayal, U., Buchmann, A., and McCarthy, D. (1988). Rules are objects too: A knowledge model for an active object-oriented database system. In *Advances in Object-Oriented Database Systems. Proc. of the 2nd Int. Workshop on Object-Oriented Database Systems*, pages 129–143. 606

Dayal, U. and Hwang, H. (1984). View definition and generalization for database integration in multibase: A system for heterogeneous distributed database. *IEEE Trans. Softw. Eng.*, SE-10(6):628–644. 147, 160, 331

Dayal, U., M.Hsu, and Ladin, R. (1991). A transactional model for long-running activities. In *Proc. 17th Int. Conf. on Very Large Data Bases*, pages 113–122. 354, 355

Dean, J. and Ghemawat, S. (2004). MapReduce: Simplified data processing on large clusters. In *Proc. 6th USENIX Symp. on Operating System Design and Implementation*, pages 137–150. 758, 763

Dean, J. and Ghemawat, S. (2010). MapReduce: a flexible data processing tool. *Commun. ACM*, 53(1):72–77. 760, 763

Demaine, E., Lopez-Ortiz, A., and Munro, J. I. (2002). Frequency estimation of internet packet streams with limited space. In *Proc. 10th Annual European Symp. on Algorithms*, pages 348–360. 743

Demers, A., Gehrke, J., Hong, M., Riedewald, M., and White, W. (2006). Towards expressive publish/subscribe systems. In *Advances in Database Technology, Proc. 10th Int. Conf. on Extending Database Technology*, pages 627–644. 741

Demers, A. J., Greene, D. H., Hauser, C., Irish, W., Larson, J., Shenker, S., Sturgis, H. E., Swinehart, D. C., and Terry, D. B. (1987). Epidemic algorithms for replicated database maintenance. In *Proc. ACM SIGACT-SIGOPS 6th Symp. on the Principles of Distributed Computing*, pages 1–12. 617

Denning, P. J. (1968). he working set model for program behavior. *Commun. ACM*, 11(5):323–333. 415

Denning, P. J. (1980). Working sets: Past and present. *IEEE Trans. Softw. Eng.*, SE-6(1):64–84. 415

Denny, M. and Franklin, M. (2005). Predicate result range caching for continuous queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 646–657. 740

Deshpande, A. and Hellerstein, J. (2004). Lifting the burden of history from adaptive query processing. In *Proc. 30th Int. Conf. on Very Large Data Bases*, pages 948–959. 739

Devine, R. (1993). Design and implementation of DDH: A distributed dynamic hashing algorithm. In *Proc. 4th Int. Conf. on Foundations of Data Organization and Algorithms*, pages 101–114. 618

DeWitt, D., Naughton, J., Schneider, D., and Seshadri, S. (1992). Practical skew handling in parallel joins. In *Proc. 22th Int. Conf. on Very Large Data Bases*, pages 27–40. 529, 548

DeWitt, D. J., Futtersack, P., Maier, D., and Velez, F. (1990). A study of three alternative workstation-server architectures for object-oriented database systems. In *Proc. 16th Int. Conf. on Very Large Data Bases*, pages 107–12. 568

DeWitt, D. J. and Gerber, R. (1985). Multi processor hash-based join algorithms. In *Proc. 11th Int. Conf. on Very Large Data Bases*, pages 151–164. 518

DeWitt, D. J., Gerber, R. H., Graek, G., Heytens, M. L., Kumar, K. B., and Muralikrishna, M. (1986). Gamma: A high performance dataflow database machine. In *Proc. 12th Int. Conf. on Very Large Data Bases*, pages 228–237. 505, 548

DeWitt, D. J. and Gray, J. (1992). Parallel database systems: The future of high performance database systems. *Commun. ACM*, 35(6):85–98. 497, 500

Dhamankar, R., Lee, Y., Doan, A., Halevy, A. Y., and Domingos, P. (2004). iMAP: Discovering complex mappings between database schemas. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 383–394. 147

Dickman, P. (1991). *Distributed Object Management in a Non-Small Graph of Autonomous Networks With Few Failures*. Ph.D. thesis, University of Cambridge, England. 581

Dickman, P. (1994). The bellerophon project: A scalable object-support architecture suitable for a large oodbms? In Özsu et al. [1994a], pages 287–299. 577

Diffie, W. and Hellman, M. E. (1976). New directions in cryptography. *IEEE Trans. Information Theory*, IT–22(6):644–654. 180

Ding, Q., Ding, Q., and Perrizo, W. (2002). Decision tree classification of spatial data streams using peano count trees. In *Proc. 2002 ACM Symp. on Applied Computing*, pages 413–417. 743

Do, H. H. and Rahm, E. (2002). COMA - A system for flexible combination of schema matching approaches. In *Proc. 28th Int. Conf. on Very Large Data Bases*, pages 610–621. 134, 142, 144, 160

Doan, A., Domingos, P., and Halevy, A. Y. (2001). Reconciling schemas of disparate data sources: A machine-learning approach. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 509–520. 145, 147

Doan, A., Domingos, P., and Halevy, A. Y. (2003a). Learning to match the schemas of data sources: A multistrategy approach. *Machine Learning*, 50(3):279–301. 145, 146, 147

Doan, A., Halevy, A., and Ives, Z. (2010). *Principles of Data Integration*. (in preparation). 159, 160

Doan, A. and Halevy, A. Y. (2005). Semantic integration research in the database community: A brief survey. *AI Magazine*, 26(1):83–94. 160

Doan, A., Madhavan, J., Dhamankar, R., Domingos, P., and Halevy, A. Y. (2003b). Learning to match ontologies on the semantic web. *VLDB J.*, 12(4):303–319. 626

Dobra, A., Garofalakis, M., Gehrke, J., and Rastogi, R. (2004). Sketch-based multi-query processing over data streams. In *Advances in Database Technology, Proc. 9th Int. Conf. on Extending Database Technology*, pages 551–568. 740

Dogac, A., Dengi, C., and Özsu, M. T. (1998a). Distributed object computing platforms. *Commun. ACM*, 41(9):95–103. 607

Dogac, A., Kalinichenko, L., Özsu, M. T., and Sheth, A., editors (1998b). *Advances in Workflow Systems and Interoperability*. Springer. 354, 359

Dogac, A., Özsu, M., Biliris, A., and Sellis, T., editors (1994). *Advances in Object-Oriented Database Systems*. Springer. 586, 607, 814

Doherty, C. and Hurley, N. (2007). Autonomic distributed data management with update accesses. In *Proc. 1st Int. Conf. on Autonomic computing and communication systems*, pages 1–8. 762

D'Oliviera, C. R. (1977). An analysis of computer decentralization. Technical Memo TM-90, Laboratory for Computer Science, Massachusetts Institute of Technology, Cambridge, Mass. 7

Dollimore, J., Nascimento, C., and Xu, W. (1994). Fine-grained object migration. In Özsu et al. [1994a], pages 182–186. 577

Domingos, P. and Hulten, G. (2000). Mining high-speed data streams. In *Proc. 6th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, pages 71–80. 743

Douglis, F., Palmer, J., Richards, E., Tao, D., Hetzlaff, W., Tracey, J., and Lin, J. (2004). Position: short object lifetimes require a delete-optimized storage system. In *Proc. 11th ACM SIGOPS European Workshop*. 726

Dowdy, L. W. and Foster, D. V. (1982). Comparative models of the file assignment problem. *ACM Comput. Surv.*, 14(2):287–313. 38, 114, 125

Draper, D., Fankhauser, P., Fernández, M., Malhotra, A., Rose, K., Rys, M., Siméon, J., and Wadler, P., editors. Xquery 1.0 and XPath 2.0 formal semantics (2007). Available from: `http://www.w3.org/TR/xquery-semantics/` [Last retrieved: January 2010]. 702

Du, W. and Elmagarmid, A. (1989). Quasi-serializability: A correctness criterion for global concurrency control in interbase. In *Proc. 15th Int. Conf. on Very Large Data Bases*, pages 347–355. 26

Du, W., Krishnamurthy, R., and Shan, M. (1992). Query optimization in a heterogeneous dbms. In *Proc. 18th Int. Conf. on Very Large Data Bases*, pages 277–291. 307, 308, 331

Du, W., Shan, M., and Dayal, U. (1995). Reducing multidatabase query response time by tree balancing. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 293–303. 287, 290, 293, 315, 331

Duschka, O. M. and Genesereth, M. R. (1997). Answering recursive queries using views. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 109–116. 160, 305, 331

Dwork, C. and Skeen, D. (1983). The inherent cost of nonblocking commitment. In *Proc. ACM SIGACT-SIGOPS 2nd Symp. on the Principles of Distributed Computing*, pages 1–11. 455

Eager, D. L. and Sevcik, K. C. (1983). Achieving robustness in distributed database systems. *ACM Trans. Database Syst.*, 8(3):354–381. 456, 493

Edwards, J., McCurley, K., and Tomlin, J. (2001). An adaptive model for optimizing performance of an incremental web crawler. In *Proc. 10th Int. World Wide Web Conf.* 666

Effelsberg, W. and Härder, T. (1984). Principles of database buffer management. *ACM Trans. Database Syst.*, 9(4):560–595. 415

Eich, M. H. (1989). Main memory database research directions. In *Int. Workshop on Database Machines*, pages 251–268. 499

Eickler, A., Gerlhof, C., and Kossmann, D. (1995). A performance evaluation of oid mapping techniques. In *Proc. 21th Int. Conf. on Very Large Data Bases*, pages 18–29. 575

Eisenberg et al., 2008 (2008). Information technology – database languages – SQL – Part 14: XML-related specifications (SQL/XML). 702

Eisner, M. J. and Severance, D. G. (1976). Mathematical techniques for efficient record segmentation in large shared databases. *J. ACM*, 23(4):619–635. 98

Elmagarmid, A., Leu, Y., Litwin, W., and Rusinkiewicz, M. (1990). A multidatabase transaction model for interbase. In *Proc. 16th Int. Conf. on Very Large Data Bases*, pages 507–518. 354

Elmagarmid, A., Rusinkiewicz, M., and Sheth, A., editors (1999). *Management of Heterogeneous and Autonomous Database Systems*. Morgan Kaufmann. 160

Elmagarmid, A. K. (1986). A survey of distributed deadlock detection algorithms. *ACM SIGMOD Rec.*, 15(3):37–45. 39, 401

Elmagarmid, A. K., editor (1992). *Transaction Models for Advanced Database Applications*. Morgan Kaufmann. 359

Elmagarmid, A. K., Soundararajan, N., and Liu, M. T. (1988). A distributed deadlock detection and resolution algorithm and its correctness proof. *IEEE Trans. Softw. Eng.*, 14(10):1443–1452. 401

Elmasri, R., Larson, J., and Navathe, S. B. (1987). Integration algorithms for database and logical database design. Technical report, Honeywell Corporate Research Center, Golden Valley, Minn. 149

Elmasri, R. and Navathe, S. B. (2011). *Fundamentals of Database Systems*. Pearson, 6 edition. 70

Embley, D. W., Jackman, D., and Xu, L. (2001). Multifaceted exploitation of metadata for attribute match discovery in information integration. In *Proc. Workshop on Information Integration on the Web*, pages 110–117. 146

Embley, D. W., Jackman, D., and Xu, L. (2002). Attribute match discovery in information integration: exploiting multiple facets of metadata. *Journal of the Brazilian Computing Society*, 8(2):32–43. 146

Epstein, R. and Stonebraker, M. (1980). Analysis of distributed data base processing strategies. In *Proc. 5th Int. Conf. on Very Data Bases*, pages 92–101. 293

Epstein, R., Stonebraker, M., and Wong, E. (1978). Query processing in a distributed relational database system. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 169–180. 209, 254, 274, 276, 292

Eswaran, K. P. (1974). Placement of records in a file and file allocation in a computer network. In *Information Processing '74*, pages 304–307. 115, 125

Eswaran, K. P., Gray, J. N., Lorie, R. A., and Traiger, I. L. (1976). The notions of consistency and predicate locks in a database system. *Commun. ACM*, 19(11):624–633. 341, 370

Evrendilek, C., Dogac, A., Nural, S., and Ozcan, F. (1997). Multidatabase query optimization. *Distrib. Parall. Databases*, 5(1):77–114. 287, 293, 316

Ezeife, C. I. and Barker, K. (1995). A comprehensive approach to horizontal class fragmentation in a distributed object based system. *Distrib. Parall. Databases*, 3(3):247–272. 563, 564, 607

Ezeife, C. I. and Barker, K. (1998). Distributed object based design: Vertical fragmentation of classes. *Distrib. Parall. Databases*, 6(4):327–360. 563

Fagin, R. (1977). Multivalued dependencies and a new normal form for relational databases. *ACM Trans. Database Syst.*, 2(3):262–278. 44

Fagin, R. (1979). Normal forms and relational database operators. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 153–160. 44

Fagin, R. (1999). Combining fuzzy information from multiple systems. *Journal of Computer and System Sciences*, 58(1):83–99. 629

Fagin, R. (2002). Combining fuzzy information: an overview. *ACM SIGMOD Rec.*, 31(2):109–118. 147

Fagin, R., Kolaitis, P. G., Miller, R. J., and Popa, L. (2005). Data exchange: semantics and query answering. *TCS*, 336(1):89–124. 159

Fagin, R., Lotem, J., and Naor, M. (2003). Optimal aggregation algorithms for middleware. *Journal of Computer and System Sciences*, 66(4):614–656. 629, 654

Fagin, R. and Vardi, M. Y. (1984). The theory of data dependencies: A survey. Research Report RJ 4321 (47149), IBM Research Laboratory, San Jose, Calif. 189

Faloutsos, C. and Christodoulakis, S. (1984). Signature files: an access method for documents and its analytical performance evaluation. *ACM Trans. Information Syst.*, 2(4):267–288. 667

Fan, W. (2004). Systematic data selection to mine concept-drifting data streams. In *Proc. 10th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, pages 128–137. 743

Fang, D., Hammer, J., and McLeod, D. (1994). An approach to behavior sharing in federated database systems. In Özsu et al. [1994a], pages 334–346. 565

Farrag, A. (1986). *Concurrency and Consistency in Database Systems*. Ph.D. thesis, Department of Computing Science, University of Alberta, Edmonton, Canada. 359

Farrag, A. A. and Özsu, M. T. (1985). A general concurrency control for database systems. In *Proc. National Computer Conf*, pages 567–573. 400

Farrag, A. A. and Özsu, M. T. (1987). Towards a general concurrency control algorithm for database systems. *IEEE Trans. Softw. Eng.*, 13(10):1073–1079. 400

Farrag, A. A. and Özsu, M. T. (1989). Using semantic knowledge of transactions to increase concurrency. *ACM Trans. Database Syst.*, 14(4):503–525. 395, 401

Fekete, A., Lynch, N., Merritt, M., and Weihl, W. (1987a). Nested transactions and read/write locking. Technical Memo MIT/LCS/TM–324, Massachusetts Institute of Technology, Cambridge, Mass. 401

Fekete, A., Lynch, N., Merritt, M., and Weihl, W. (1987b). Nested transactions, conflict-based locking, and dynamic atomicity. Technical Memo MIT/LCS/TM–340, Massachusetts Institute of Technology, Cambridge, Mass. 401

Fekete, A., Lynch, N., Merritt, M., and Weihl, W. (1989). Commutativity-based locking for nested transactions. Technical Memo MIT/LCS/TM-370b, Massachusetts Institute of Technology, Cambridge, Mass. 401, 594

Fernandez, E. B., Summers, R. C., and Wood, C. (1981). *Database Security and Integrity*. Addison Wesley. 180

Fernandez, M., Florescu, D., and Levy, A. (1997). A query language for a web-site management system. *ACM SIGMOD Rec.*, 26(3):4–11. 676

Fernández, M. F., Siméon, J., Choi, B., Marian, A., and Sur, G. (2003). Implementing XQuery 1.0: The Galax experience. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 1077–1080. 698, 702

Ferreira, P. and Shapiro, M. (1994). Garbage collection and dsm consistency. In *Proc. of the First Symposium on Operating Systems Design and Implementation*, pages 229–241. 581

Fessant, F. L., Piumarta, I., and Shapiro, M. (1998). An implementation of complete, asynchronous, distributed garbage collection. In *Proc. ACM SIGPLAN Conf. on Programming Language Design and Implementation*, pages 152–161. 582

Fiebig, T., Helmer, S., Kanne, C.-C., Moerkotte, G., Neumann, J., Schiele, R., and Westmann, T. (2002). Anatomy of a native XML base management system. *VLDB J.*, 11(4):292–314. 699

Fisher, M. K. and Hochbaum, D. S. (1980). Database location in computer networks. *J. ACM*, 27(4):718–735. 121

Fisher, P. S., Hollist, P., and Slonim, J. (1980). A design methodology for distributed data bases. In *Digest of Papers – COMPCON*, pages 199–202. 125

Florentin, J. J. (1974). Consistency auditing of databases. *Comp. J.*, 17(1):52–58. 188, 202

Florescu, D., Koller, D., and Levy, A. (1997). Using probabilistic information in data integration. In *Proc. 23th Int. Conf. on Very Large Data Bases*, pages 216–225. 564

Florescu, D., Levy, A., and Mendelzon, A. (1998). Database techniques for the World-Wide Web: a survey. *ACM SIGMOD Rec.*, 27(3):59–74. 657, 676

Folkert, N., Gupta, A., Witkowski, A., Subramanian, S., Bellamkonda, S., Shankar, S., Bozkaya, T., and Sheng, L. (2005). Optimizing refresh of a set of materialized views. In *Proc. 31st Int. Conf. on Very Large Data Bases*, pages 1043–1054. 738

Foster, D. V. and Browne, J. C. (1976). File assignment in memory hierarchies. In Gelenbe, I. E., editor, *Modelling and Performance Evaluation of Computer Systems*, pages 119–127. North-Holland. 125

Franklin, M., Livny, M., and Carey, M. (1997). Transactional client-server cache consistency: Alternatives and performance. *ACM Trans. Database Syst.*, 22(3):315–367. 572

Franklin, M. J., Carey, M., and Livny, M. (1992). Global memory management in client-server dbms architectures. In *Proc. 18th Int. Conf. on Very Large Data Bases*, pages 596–609. 210, 571

Franklin, M. J. and Carey, M. J. (1994). Client-server caching revisited. In Özsu et al. [1994a], pages 57–78. 572, 573

Franklin, M. J., Jonsson, B. T., and Kossmann, D. (1996). Performance tradeoffs for client-server query processing. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 149–160. 214

Freeley, M., Morgan, W., and Pighin, F. (1995). Implementing global memory management in a workstation cluster. In *Proc. 15th ACM Symp. on Operating Syst. Principles*, pages 201–212. 210

Freytag, J. C. (1987). A rule-based view of query optimization. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 173–180. 583

Freytag, J. C., Maier, D., and Vossen, G. (1994). *Query Processing for Advanced Database Systems*. Morgan Kaufmann. 220

Friedman, M., Levy, A. Y., and Millstein, T. D. (1999). Navigational plans for data integration. In *Proc. 16th National Conf on Artificial Intelligence and 11th Innovative Applications of Artificial Intelligence Conf.*, pages 67–73. 133

Fung, C. W., Karlaplem, K., and Li, Q. (1996). An analytical approach towards evaluating method induced vertical partitioning algorithms. Technical Report HKUST96-33, Department of Computer Science, Hong Kong University of Science and Technology. 564

Furtado, C., Lima, A., Pacitti, E., Valduriez, P., and Mattoso, M. (2005). Physical and virtual partitioning in olap database clusters. In *Proc. Int. Symp. Computer Architecture and High Performance Computing*, pages 143–150. 544, 548

Furtado, C., Lima, A., Pacitti, E., Valduriez, P., and Mattoso, M. (2006). Adaptive hybrid partitioning for olap query processing in a database cluster. *Int. J. High Perf. Comput. and Networking*. To appear. 544, 548

Fushimi, S., Kitsuregawa, M., and Tanaka, H. (1986). An overview of the system software of a parallel relational database machine grace. In *Proc. 12th Int. Conf. on Very Large Data Bases*, pages 209–219. 505

Gaber, M., Zaslavsky, A., and Krishnaswamy, S. (2005). Mining data streams: A review. *ACM SIGMOD Rec.*, 34(2):18–26. 742, 762

Galhardas, H., Florescu, D., Shasha, D., Simon, E., and Saita, C.-A. (2001). Declarative data cleaning: Language, model, and algorithms. In *Proc. 27th Int. Conf. on Very Large Data Bases*, pages 371–380. 158

Gallaire, H., Minker, J., and Nicolas, J.-M. (1984). Logic and databases: A deductive approach. *ACM Comput. Surv.*, 16(2):153–186. 47

Gama, J., Medas, P., and Rodrigues, P. (2005). Learning decision trees from dynamic data streams. In *Proc. 2005 ACM Symp. on Applied Computing*, pages 573–577. 743

Gançarski, S., Naacke, H., Pacitti, E., and Valduriez, P. (2002). Parallel processing with autonomous databases in a cluster system. In *Proc. Int. Conf. on Cooperative Information Systems*, pages 410–428. 540, 548

Gançarski, S., Naacke, H., Pacitti, E., and Valduriez, P. (2007). The leganet system: Freshness-aware transaction routing in a database cluster. *Inf. Syst.*, 32(7):320–343. 541, 548

Ganesan, P., Yang, B., and Garcia-Molina, H. (2004). One torus to rule them all: Multidimensional queries in p2p systems. In *Proc. 7th Int. Workshop on the World Wide Web and Databases*, pages 19–24. 622

Ganti, Gehrke, and Ramakrishnan (2002). Mining data streams under block evolution. *SIGKDD Explorations*, pages 1–10. 743

Gao, S., Sperberg-McQueen, C. M., and Thompson, H. S., editors. W3C XML schema definition language (XSD) 1.1 part 1: Structures (2009). Available from: `http://www.w3.org/TR/xmlschema11-1/` [Last retrieved: January 2010]. 693

Garcia-Molina, H. (1979). *Performance of Update Algorithms for Replicated Data in a Distributed Database*. Ph.D. thesis, Department of Computer Science, Stanford University, Stanford, Calif. 390, 401

Garcia-Molina, H. (1982). Elections in distributed computing systems. *IEEE Trans. Comput.*, C-31(1):48–59. 440

Garcia-Molina, H. (1983). Using semantic knowledge for transaction processing in a distributed database. *ACM Trans. Database Syst.*, 8(2):186–213. 352, 395, 401

Garcia-Molina, H., Gawlick, D., Klein, J., Kleissner, K., and Salem, K. (1990). Coordinating multi-transaction activities. Technical Report CS-TR-247-90, Department of Computer Science, Princeton University. 352, 353, 397

Garcia-Molina, H., Papakonstantinou, Y., Quass, D., Rajaraman, A., Sagiv, Y., Ullman, J. D., Vassalos, V., and Widom, J. (1997). The TSIMMIS approach to mediation: Data models and languages. *J. Intell. Information Syst.*, 8(2):117–132. 160

Garcia-Molina, H. and Salem, K. (1987). Sagas. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 249–259. 351, 352, 397

Garcia-Molina, H., Ullman, J. D., and Widom, J. (2002). *Database Systems – The Complete Book*. Prentice-Hall. 70

Garcia-Molina, H. and Wiederhold, G. (1982). Read–only transactions in a distributed database. *ACM Trans. Database Syst.*, 7(2):209–234. 401

Garofalakis, M. N. and Ioannidis, Y. E. (1996). Multi-dimensional resource scheduling for parallel queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 365–376. 530, 548

Garza, J. F. and Kim, W. (1988). Transaction management in an object-oriented database system. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 37–45. 597, 600

Gastonian, R. (1983). The auragen system 4000. *Q. Bull. IEEE TC on Data Eng.*, 6(2). 456

Gavish, B. and Pirkul, H. (1986). Computer and database location in distributed computer systems. *IEEE Trans. Comput.*, C-35(7):583–590. 125

GE (1976). *MADMAN User Manual*. General Electric Company, Schenectady, N.Y. 390

Gedik, B., Wu, K.-L., Yu, P. S., and Liu, L. (2005). Adaptive load shedding for windowed stream joins. In *Proc. 14th ACM Int. Conf. on Information and Knowledge Management*, pages 171–178. 740

Gelenbe, E. and Gardy, D. (1982). The size of projections of relations satisfying a functional dependency. In *Proc. 8th Int. Conf. on Very Data Bases*, pages 325–333. 254

Gelenbe, E. and Sevcik, K. (1978). Analysis of update synchronization for multiple copy databases. In *Proc. 3rd Berkeley Workskop on Distributed Data Management and Computer Networks*, pages 69–88. 401

Georgakopoulos, D., Hornick, M., and Sheth, A. (1995). An overview of workflow management: From process modeling to workflow automation infrastructure. *Distrib. Parall. Databases*, 3:119–153. 354, 359

Gerlhof, C. and Kemper, A. (1994). A multi-threaded architecture for prefetching in object bases. In Jarke, M., Jr., J. A. B., and Jeffery, K. G., editors, *Advances in Database Technology, Proc. 4th Int. Conf. on Extending Database Technology*, volume 779 of *Lecture Notes in Computer Science*, pages 351–364. Springer. 568

Ghanem, T., Aref, W., and Elmagarmid, A. (2006). Exploiting predicate-window semantics over data streams. *ACM SIGMOD Rec.*, 35(1):3–8. 727

Ghemawat, S. (1995). *The Modified Object Buffer: A Storage Management Technique for Object-Oriented Databases*. Ph.D dissertation, Massachusetts Institute of Technology, Cambridge, Mass. 571

Ghemawat, S., Gobioff, H., and Leung, S.-T. (2003). The Google file system. In *Proc. 19th ACM Symp. on Operating System Principles*, pages 29–43. 753, 763

Gibbons, P. and Tirthapura, S. (2002). Distributed streams algorithms for sliding windows. In *Proc. 14th ACM Symp. on Parallel Algorithms and Architectures*, pages 63–72. 737

Gibbons, T. (1976). *Integrity and Recovery in Computer Systems*. NCC Publications. 455

Gifford, D. K. (1979). Weighted voting for replicated data. In *Proc. 7th ACM Symp. on Operating System Principles*, pages 50–159. 487

Gilbert, A. C., Kotidis, Y., Muthukrishnan, S., and Strauss, M. J. (2001). Surfing wavelets on streams: One-pass summaries for approximate aggregate queries. In *Proc. 27th Int. Conf. on Very Large Data Bases*, pages 79–88. 726

Gligor, V. and Popescu-Zeletin, R. (1986). Transaction management in distributed heterogeneous database management systems. *Inf. Syst.*, 11(4):287–297. 25

Gligor, V. D. and Luckenbaugh, G. L. (1984). Interconnecting heterogeneous database management systems. *Comp.*, 17(1):33–43. 40

Golab, L. (2006). *Sliding Window Query Processing over Data Streams*. PhD thesis, University of Waterloo. 763

Golab, L., Garg, S., and Özsu, M. T. (2004). On indexing sliding windows over on-line data streams. In *Advances in Database Technology, Proc. 9th Int. Conf. on Extending Database Technology*, pages 712–729. 736

Golab, L., Johnson, T., Seidel, J. S., and Shkapenyuk, V. (2009). Stream warehousing with DataDepot. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 847–854. 761

Golab, L. and Özsu, M. T. (2003a). Issues in data stream management. *ACM SIGMOD Rec.*, 32(2):5–14. 762, 763

Golab, L. and Özsu, M. T. (2003b). Processing sliding window multi-joins in continuous queries over data streams. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 500–511. 733

Golab, L. and Özsu, M. T. (2010). *Data Stream Systems*. Morgan & Claypool. 761, 762, 763

Goldberg, A. and Robson, D. (1983). *SmallTalk-80: The Language and Its Implementation*. Addison Wesley. 559

Goldman, K. J. (1987). Data replication in nested transaction systems. Technical Report MIT/LCS/TR-390, Massachusetts Institute of Technology, Cambridge, Mass. 401

Goldman, R. and Widom, J. (1997). Dataguides : Enabling query formulation and optimization in semistructured databases. In *Proc. 23th Int. Conf. on Very Large Data Bases*, pages 436–445. 675, 701

Gonnet, G. H. and Tompa, F. W. (1987). Mind your grammar: A new approach to modelling text. In *Proc. 13th Int. Conf. on Very Large Data Bases*, pages 339–346. 690

Goodman, J. R. and Woest, P. J. (1988). The wisconsin multicube: A new large-scale cache-coherent multiprocessor. Technical Report TR766, University of Wisconsin-Madison. 506, 548

Goodman, N., Suri, R., and Tay, Y. C. (1983). A simple analytic model for performance of exclusive locking in database systems. In *Proc. 2nd ACM SIGACT–SIGMOD Symp. on Principles of Database Systems*, pages 203–215. 401

Gottlob, G., Koch, C., and Pichler, R. (2005). Efficient algorithms for processing XPath queries. *ACM Trans. Database Syst.*, 30(2):444–491. 700

Gounaris, A., Paton, N., Fernandes, A., and Sakellariou, R. (2002a). Adaptive query processing: A survey. In *Proc. British National Conf. on Databases*, pages 11–25. 739

Gounaris, A., Paton, N. W., Fernandes, A. A. A., and Sakellariou, R. (2002b). Adaptive query processing: A survey. In *Proc. British National Conf. on Databases*, pages 11–25. 320, 321, 331

Graefe, G. (1990). Encapsulation of parallelism in the volcano query processing systems. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 102–111. 503, 548

Graefe, G. (1993). Query evaluation techniques for large databases. *ACM Comput. Surv.*, 25(2):73–170. 220, 292, 547

Graefe, G. (1994). Volcano - an extensible and parallel query evaluation system. *IEEE Trans. Knowl. and Data Eng.*, 6(1):120–135. 267

Graefe, G. and DeWitt, D. (1987). The exodus optimizer generator. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 160–172. 583

Graefe, G. and Maier, D. (1988). Query optimization in object-oriented database systems: The REVELATION project. Technical Report CS/E 88-025, Oregon Graduate Center. 583, 586

Graefe, G. and McKenna, W. (1993). The volcano optimizer generator. In *Proc. 9th Int. Conf. on Data Engineering*, pages 209–218. 320, 321, 586

Grant, J. (1984). Constraint preserving and lossless database transformations. *Inf. Syst.*, 9(2):139–146. 79

Grapa, E. and Belford, G. G. (1977). Some theorems to aid in solving the file allocation problem. *Commun. ACM*, 20(11):878–882. 125

Gravano, L., Garcia-Molina, H., and Tomasic, A. (1999). Gloss: Text-source discovery over the internet. *ACM Trans. Database Syst.*, 24(2):229–264. 689

Gray, J. (1981). The transaction concept: Virtues and limitations. In *Proc. 7th Int. Conf. on Very Data Bases*, pages 144–154. 337

Gray, J. (1985). Why do computers stop and what can be done about it. Technical Report 85-7, Tandem Computers, Cupertino, Calif. 455, 456

Gray, J. (1987). Why do computers stop and what can be done about it. In *CIPS (Canadian Information Processing Society) Edmonton '87 Conf. Tutorial Notes*, Edmonton, Canada. 350, 410

Gray, J. (1989). Transparency in its place – the case against transparent access to geographically distributed data. Technical Report TR89.1, Tandem Computers Inc, Cupertino, Calif. 11

Gray, J., Helland, P., O'Neil, P. E., and Shasha, D. (1996). The dangers of replication and a solution. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 173–182. 460, 493

Gray, J. and Reuter, A. (1993). *Transaction Processing: Concepts and Techniques*. Morgan Kaufmann. 358, 396, 401

Gray, J. N. (1979). Notes on data base operating systems. In Bayer, R., Graham, R. M., and Seegmüller, G., editors, *Operating Systems: An Advanced Course*, pages 393–481. Springer. 39, 359, 419, 425, 426, 431, 456

Gray, J. N., Lorie, R. A., Putzolu, G. R., and Traiger, I. L. (1976). Granularity of locks and degrees of consistency in a shared data base. In Nijssen, G. M., editor, *Modelling in Data Base Management Systems*, pages 365–394. North-Holland. 345

Gray, J. N., McJones, P., Blasgen, M., Lindsay, B., Lorie, R., Price, T., Putzolu, F., and Traiger, I. (1981). The recovery manager of the system r database manager. *ACM Comput. Surv.*, 13(2):223–242. 411, 419, 426, 456

Grefen, P. and Widom, J. (1997). Protocols for integrity constraint checking in federated databases. *Distrib. Parall. Databases*, 5(4):327–355. 200, 202

Griffiths, P. P. and Wade, B. W. (1976). An authorization mechanism for a relational database system. *ACM Trans. Database Syst.*, 1(3):242–255. 182, 201

Grossman, R. L. and Gu, Y. (2009). On the varieties of clouds for data intensive computing. *Q. Bull. IEEE TC on Data Eng.*, 32(1):44–50. 745

Group, E. D. S. E. D. (1990). Eds-collaborating for a high-performance parallel relational database. In *Proc. ESPRIT Conf*, pages 274–295. 505, 548

Gruber, O. and Amsaleg, L. (1994). Object grouping in eos. In Özsu et al. [1994a], pages 117–131. 579

Grust, T., van Keulen, M., and Teubner, J. (2003). Staircase join: Teach a relational dbms to watch its (axis) steps. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 524–525. 700

Gudgin, M., Hadley, M., Mendelsohn, N., Moreau, J.-J., Nielsen, H. F., Karmarkar, A., and Lafon, Y., editors. Simple object protocol (SOAP) version 1.2 (2007). Available from: http://www.w3.org/TR/soap12 [Last retrieved: December 2009]. 690

Guerrini, G., Bertino, E., and Bal, R. (1998). A formal definition of the chimera object-oriented data model. *J. Intell. Information Syst.*, 11(1):5–40. 607

Guha, S. and McGregor, A. (2006). Approximate quantiles and the order of the stream. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 273–279. 725

Guha, S., Meyerson, A., Mishra, N., and Motwani, R. (2003). Clustering data streams: Theory and practice. *IEEE Trans. Knowl. and Data Eng.*, 15(3):515–528. 743

Gulisano, V., Jimenez-Peris, R., Patino-Martinez, M., and Valduriez, P. (2010). StreamCloud: A large scale data streaming system. In *Proc. 30th Int. Conf. on Distributed Computing Systems*. 762

Gulli, A. and Signorini, A. (2005). The indexable web is more than 11.5 billion pages. In *Proc. 14th Int. World Wide Web Conf.*, pages 902 – 903. 657

Gummadi, P. K., Gummadi, R., Gribble, S. D., Ratnasamy, S., Shenker, S., and Stoica, I. (2003). The impact of DHT routing geometry on resilience and proximity. In *Proc. ACM Int. Conf. on Data Communication*, pages 381–394. 619

Güntzer, U., Kießling, W., and Balke, W.-T. (2000). Optimizing multi-feature queries for image databases. In *Proc. 26th Int. Conf. on Very Large Data Bases*, pages 419–428. 629, 654

Guo, H., Larson, P.-A., Ramakrishnan, R., and Goldstein, J. (2004). Relaxed currency and consistency: How to say "good enough" in sql. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 815–826. 540

Gupta, A., Agrawal, D., and Abbadi, A. E. (2003). Approximate range selection queries in peer-to-peer systems. In *Proc. 1st Biennial Conf. on Innovative Data Systems Research*, pages 141–151. 642

Gupta, A., Jagadish, H., and Mumick, I. S. (1996). Data integration using self-maintainable views. In *Advances in Database Technology, Proc. 5th Int. Conf. on Extending Database Technology*, pages 140–144. 179, 180

Gupta, A. and Mumick, I. S. (1999a). Maintenance of materialized views: Problems, techniques, and applications. In Gupta and Mumick [1999c], chapter 11, pages 145–156. 178, 201

Gupta, A. and Mumick, I. S., editors (1999b). *Materialized Views: Techniques, Implementations, and Applications*. M.I.T. Press. 132

Gupta, A. and Mumick, I. S., editors (1999c). *Materialized Views: Techniques, Implementations, and Applications*. M.I.T. Press. 176, 201, 794

Gupta, A., Mumick, I. S., and Subrahmanian, V. S. (1993). Maintaining views incrementally. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 157–166. 179, 201

Haas, L. (2007). Beauty and the beast: The theory and practice of information integration. In *Proc. 11th Int. Conf. on Database Theory*, pages 28–43. 160

Haas, L., Kossmann, D., Wimmers, E., and Yang, J. (1997a). Optimizing queries across diverse data sources. In *Proc. 23th Int. Conf. on Very Large Data Bases*, pages 276–285. 317, 331

Haas, L. M., Kossmann, D., Wimmers, E. L., and Yang, J. (1997b). Optimizing queries across diverse data sources. In *Proc. 23th Int. Conf. on Very Large Data Bases*, pages 276–285. 160

Haas, P. and Hellerstein, J. (1999a). Ripple joins for online aggregation. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 287–298. 732

Haas, P. J. and Hellerstein, J. M. (1999b). Ripple joins for online aggregation. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 287–298. 322, 325, 331

Haderle, C. M. D., Lindsay, B., Pirahesh, H., and Schwarz, P. (1992). Aries: A transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging. *ACM Trans. Database Syst.*, 17(1):94–162. 401, 418

Hadzilacos, T. and Hadzilacos, V. (1991). Transaction synchroniation in object bases. *J. Comp. and System Sci.*, 43(1):2–24. 597

Hadzilacos, V. (1988). A theory of reliability in database systems. *J. ACM*, 35(1):121–145. 429, 456, 596

Haessig, K. and Jenny, C. J. (1980). An algorithm for allocating computational objects in distributed computing systems. Research Report RZ 1016, IBM Research Laboratory, Zurich. 125

Halatchev, M. and Gruenwald, L. (2005). Estimating missing values in related sensor data streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 83–94. 744

Halevy, A., Rajaraman, A., and Ordille, J. (2006). Data integration: The teenage years. In *Proc. 32nd Int. Conf. on Very Large Data Bases*, pages 9–16. 160

Halevy, A. Y. (2001). Answering queries using views: A survey. *VLDB J.*, 10(4):270–294. 301, 304, 331

Halevy, A. Y., Ashish, N., Bitton, D., Carey, M., Draper, D., Pollock, J., Rosenthal, A., and Sikka, V. (2005). Enterprise information integration: Successes, challenges and controversies. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 778–787. 131

Halevy, A. Y., Etzioni, O., Doan, A., Ives, Z. G., Madhavan, J., McDowell, L., and Tatarinov, I. (2003). Crossing the structure chasm. In *Proc. 1st Biennial Conf. on Innovative Data Systems Research*. 159

Halici, U. and Dogac, A. (1989). Concurrency control in distributed databases through time intervals and short-term locks. *IEEE Trans. Softw. Eng.*, 15(8):994–995. 401

Hammad, M., Aref, W., and Elmagarmid, A. (2003a). Stream window join: Tracking moving objects in sensor-network databases. In *Proc. 15th Int. Conf. on Scientific and Statistical Database Management*, pages 75–84. 733

Hammad, M., Aref, W., and Elmagarmid, A. (2005). Optimizing in-order execution of continuous queries over streamed sensor data. In *Proc. 17th Int. Conf. on Scientific and Statistical Database Management*, pages 143–146. 733

Hammad, M., Aref, W., Franklin, M., Mokbel, M., and Elmagarmid, A. (2003b). Efficient execution of sliding window queries over data streams. Technical Report CSD TR 03-035, Purdue University. 733, 734, 735, 736

Hammad, M., Mokbel, M., Ali, M., Aref, W., Catlin, A., Elmagarmid, A., Eltabakh, M., Elfeky, M., Ghanem, T., Gwadera, R., Ilyas, I., Marzouk, M., and Xiong, X. (2004). Nile: a query processing engine for data streams. In *Proc. 20th Int. Conf. on Data Engineering*, page 851. 735, 736

Hammer, M. and Niamir, B. (1979). A heuristic approach to attribute partitioning. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 93–101. 99, 125

Hammer, M. and Shipman, D. W. (1980). Reliability mechanisms for sdd-1: A system for distributed databases. *ACM Trans. Database Syst.*, 5(4):431–466. 440, 456

Han, D., Xiao, C., Zhou, R., Wang, G., Huo, H., and Hui, X. (2006). Load shedding for window joins over streams. In *Proc. 7th Int. Conf. on Web-Age Information Management:*, pages 472–483. 740

Hanson, E., Carnes, C., Huang, L., Konyala, M., and Noronha, L. (1999). Scalable trigger processing. In *Proc. 15th Int. Conf. on Data Engineering*, pages 266–275. 741

Härder, T. and Reuter, A. (1983). Principles of transaction-oriented database recovery. *ACM Comput. Surv.*, 15(4):287–317. 39, 411, 413, 420, 421, 423, 424, 456

Harizopoulos, S., Shah, M. A., Meza, J., and Ranganathan, P. (2009). Energy efficiency: The new holy grail of data management systems research. In *Proc. 4th Biennial Conf. on Innovative Data Systems Research*. 762

Harvey, N. J. A., Jones, M. B., Saroiu, S., Theimer, M., and Wolman, A. (2003). SkipNet: A scalable overlay network with practical locality properties. In *Proc. 4th USENIX Symp. on Internet Tech. and Systems*. 618, 622, 642

He, B., Chang, K. C.-C., and Han, J. (2004). Mining complex matchings across web query interfaces. In *Proc. ACM SIGMOD Workshop on Research Issues in Data Mining and Knowledge Discovery*, pages 3–10. 149

He, Q. and Ling, T. W. (2006). An ontology-based approach to the integration of entity-relationship schemas. *Data & Knowl. Eng.*, 58(3):299–326. 134

Hedley, Y. L., Younas, M., James, A., and Sanderson, M. (2004a). A two-phase sampling technique for information extraction from hidden web databases. In *WIDM04*, pages 1–8. 688

Hedley, Y.-L., Younas, M., James, A. E., and Sanderson, M. (2004b). Query-related data extraction of hidden web documents. In *Proc. 30th Annual Int. ACM SIGIR Conf. on Research and Development in Information Retrieval*, pages 558–559. 687

Heimbigner, D. and McLeod, D. (1985). A federated architecture for information management. *ACM Trans. Information Syst.*, 3(3):253–278. 36

Helal, A. A., Heddaya, A. A., and Bhargava, B. B. (1997). *Replication Techniques in Distributed Systems*. Kluwer Academic Publishers. 456, 486, 493

Hellerstein, J. M., Franklin, M. J., Chandrasekaran, S., Deshpande, A., Hildrum, K., Madden, S., Raman, V., and Shah, M. A. (2000). Adaptive query processing: Technology in evolution. *Q. Bull. IEEE TC on Data Eng.*, 23(2):7–18. 320, 331

Hellerstein, J. M., Haas, P., and Wang, H. (1997). Online aggregation. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 171–182. 732

Hellerstein, J. M. and Stonebraker, M. (1993). Predicate migration: Optimizing queries with expensive predicates. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 267–276. 323

Herlihy, M. (1987). Concurrency versus availability: Atomicity mechanisms for replicated data. *ACM Trans. Comp. Syst.*, 5(3):249–274. 456, 493

Herlihy, M. (1990). Apologizing versus asking permission: Optimistic concurrency control for abstract data types. *ACM Trans. Database Syst.*, 15(1):96–124. 594, 595

Herman, D. and Verjus, J. P. (1979). An algorithm for maintaining the consistency of multiple copies. In *Proc. 1st Int. Conf. on Distributed Computing Systems*, pages 625–631. 382

Hernández, M. A. and Stolfo, S. J. (1998). Real-world data is dirty: Data cleansing and the merge/purge problem. *Proc. ACM SIGMOD Workshop on Research Issues in Data Mining and Knowledge Discovery*, 2(1):9–37. 158

Herrmann, U., Dadam, P., Küspert, K., Roman, E. A., and Schlageter, G. (1990). A lock technique for disjoint and non-disjoint complex objects. In *Advances in Database Technology, Proc. 2nd Int. Conf. on Extending Database Technology*, pages 219–237. Springer. 602

Hersh, W. (2001). Managing gigabytes - compressing and indexing documents and images (second edition). *Inf. Ret.*, 4(1):79–80. 667

Hevner, A. R. and Schneider, G. M. (1980). An integrated design system for distributed database networks. In *Digest of Papers - COMPCON*, pages 459–465. 125

Hevner, A. R. and Yao, S. B. (1979). Query processing in distributed database systems. *IEEE Trans. Softw. Eng.*, 5(3):177–182. 255

Hirate, Y., Kato, S., and Yamana, H. (2006). Web structure in 2005. In *Proc. 4th Int. Workshop on Algorithms and Models for the Web-Graph*, pages 36 – 46. 657

Hoffer, H. A. and Severance, D. G. (1975). The use of cluster analysis in physical data base design. In *Proc. 1st Int. Conf. on Very Data Bases*, pages 69–86. 99, 102, 105, 125

Hoffer, J. A. (1975). *A Clustering Approach to the Generation of Subfiles for the Design of a Computer Data Base*. Ph.D. thesis, Department of Operations Research, Cornell University, Ithaca, N.Y. 125

Hoffman, J. L. (1977). *Model Methods for Computer Security and Privacy*. Prentice-Hall. 181, 201

Hofri, M. (1994). On timeout for global deadlock detection in decentralized database systems. *Inf. Proc. Letters*, 51(6):295–302. 401

Hong, W. (1992). Exploiting inter-operation parallelism in xprs. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 19–28. 503, 530, 533, 548

Hsiao, D., editor (1983). *Advanced Database Machine Architectures*. Prentice-Hall. 498

Hsiao, H. I. and DeWitt, D. (1991). A performance study of three high-availability data replication strategies. In *Proc. Int. Conf. on Parallel and Distributed Information Systems*, pages 18–28. 511, 512

Hsu, M., editor (1993). *IEEE Quart. Bull. Data Eng., Special Issue on Workflow and Extended Transaction Systems*, volume 16. IEEE Computer Society. 354

Huebsch, R., Hellerstein, J., Lanham, N., Loo, B. T., Shenker, S., and Stoica, I. (2003). Querying the internet with pier. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 321–332. 641

Hull, R. (1997). Managing semantic heterogeneity in databases: A theoretical perspective. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 51–61. 160

Hulten, G., Spencer, L., and Domingos, P. (2001). Mining time-changing data streams. In *Proc. 7th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, pages 97–106. 743, 762

Hunt, H. B. and Rosenkrantz, D. J. (1979). The complexity of testing predicate locks. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 127–133. 233

Hwang, D. J. (1987). Constructing a highly-available location service for a distributed environment. Technical Report MIT/LCS/TR-410, Massachusetts Institute of Technology, Cambridge, Mass. 577

Ibaraki, T. and Kameda, T. (1984). On the optimal nesting order for computing *n*-relation joins. *ACM Trans. Database Syst.*, 9(3):482–502. 207, 220, 245

Ilyas, I. F., Beskales, G., and Soliman, M. A. (2008). A survey of top-k query processing techniques in relational database systems. *ACM Comput. Surv.*, 40(4):1–58. 628, 654

Inmon, W. (1992). *Building the Data Warehouse*. John Wiley & Sons. 131

Ioannidis, Y. (1996). Query optimization. In Tucker, A., editor, *The Computer Science and Engineering Handbook*, pages 1038–1054. CRC Press. 292

Ioannidis, Y. and Wong, E. (1987). Query optimization by simulated annealing. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 9–22. 212, 249, 586

Ipeirotis, P. G. and Gravano, L. (2002). Distributed search over the hidden web: Hierarchical database sampling and selection. In *Proc. 28th Int. Conf. on Very Large Data Bases*, pages 394–405. 687, 688, 689

Irani, K. B. and Khabbaz, N. G. (1982). A methodology for the design of communication networks and the distribution of data in distributed computer systems. *IEEE Trans. Comput.*, C-31(5):419–434. 125

Isloor, S. S. and Marsland, T. A. (1980). The deadlock problem: An overview. *Comp.*, 13(9):58–78. 39, 401

Jagadish, H. V., Ooi, B. C., Tan, K.-L., Vu, Q. H., and Zhang, R. (2006). Speeding up search in peer-to-peer networks with a multi-way tree structure. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 1–12. 622

Jagadish, H. V., Ooi, B. C., and Vu, Q. H. (2005). BATON: A balanced tree structure for peer-to-peer networks. In *Proc. 31st Int. Conf. on Very Large Data Bases*, pages 661–672. 622, 643

Jajodia, S., Atluri, V., Keefe, T. F., McCollum, C. D., and Mukkamala, R. (2001). Multilevel security transaction processing. *J. Computer Security*, 9(3):165–195. 187, 202

Jajodia, S. and Mutchler, D. (1987). Dynamic voting. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 227–238. 456, 493

Jajodia, S. and Sandhu, R. S. (1991). Towards a multilevel secure relational data model. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 50–59. 181, 202

Jarke, M. and Koch, J. (1984). Query optimization in database systems. *ACM Comput. Surv.*, 16(2):111–152. 211, 220, 241

Jarke, M., Lenzerini, M., Vassiliou, Y., and Vassiliadis, P. (2003). *Fundamentals of Data Warehouses*. Springer, 2 edition. 131

Jenq, B., Woelk, D., Kom, W., and Lee, W. L. (1990). Query processing in distributed orion. In *Advances in Database Technology, Proc. 2nd Int. Conf. on Extending Database Technology*, pages 169–187. Springer. 587

Jhingran, A. D., Mattos, N., and Pirahesh, H. (2002). Information integration: A research agenda. *IBM Systems J.*, 41(4):555–562. 131

Jiang, H., Lu, H., 0011, W. W., and Ooi, B. C. (2003). Xr-tree: Indexing XML data for efficient structural joins. In *Proc. 19th Int. Conf. on Data Engineering*, pages 253–263. 701

Jiang, N. and Gruenwald, L. (2006). Research issues in data stream association rule mining. *ACM SIGMOD Rec.*, 35(1):14–19. 743

Jiang, Q. and Chakravarthy, S. (2004). Scheduling strategies for processing continuous queries over streams. In *Proc. British National Conf. on Databases*, pages 16–30. 735

Jiménez-Peris, R., Patiño-Martínez, M., and Alonso, G. (2002). Non-intrusive, parallel recovery of replicated data. In *Proc. 21st Symp. on Reliable Distributed Systems*, pages 150–159. 546, 548

Jiménez-Peris, R., Patiño-Martínez, M., Alonso, G., and Kemme, B. (2003). Are quorums an alternative for data replication? *ACM Trans. Database Syst.*, 28(3):257–294. 489, 548

Jiménez-Peris, R., Patiño-Martínez, M., and Kemme, B. (2007). Enterprise grids: Challenges ahead. *J. Grid Comp.*, 5(3):283–294. 748

Jiménez-Peris, R., Patiño-Martínez, M., Kemme, B., and Alonso, G. (2002). Improving the scalability of fault-tolerant database clusters. In *Proc. 22nd Int. Conf. on Distributed Computing Systems*, pages 477–484. 482, 491, 548

Jones, A. K. (1979). The object model: A conceptual tool for structuring software. In Bayer, R., Graham, R. M., and Seegmüller, G., editors, *Operating Systems: An Advanced Course*, pages 7–1. Springer. 555

Josifovski, V., Fontoura, M., and Barta, A. (2005). Querying XML streams. *VLDB J.*, 14(2):197–210. 700

Jr, A. M. J. and Malek, M. (1988). Survey of software tools for evaluating reliability, availability and serviceability. *ACM Comput. Surv.*, 20(4):227–269. 455

Kabra, N. and DeWitt, D. J. (1998). Efficient mid-query re-optimization of sub-optimal query execution plans. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 106–117. 739

Kaelbling, L. P., Littman, M. L., and Moore, A. P. (1996). Reinforcement learning: A survey. *J. Artificial Intel. Res.*, 4:237–285. 666

Kaiser, G. (1989). Transactions for concurrent object-oriented programming systems. In *Proc. ACM SIGPLAN Workshop on Object-Based Concurrent Programming*, pages 136–138. 593

Kalogeraki, V., Gunopulos, D., and Zeinalipour-Yazti, D. (2002). A local search mechanism for peer-to-peer networks. In *Proc. 11th Int. Conf. on Information and Knowledge Management*, pages 300–307. 617

Kambayashi, Y., Yoshikawa, M., and Yajima, S. (1982). Query processing for distributed databases using generalized semi–joins. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 151–160. 272, 292

Kang, J., Naughton, J., and Viglas, S. (2003). Evaluating window joins over unbounded streams. In *Proc. 19th Int. Conf. on Data Engineering*, pages 341–352. 733, 738

Kanne, C.-C. and Moerkotte, G. (2000). Efficient storage of XML data. In *Proc. 16th Int. Conf. on Data Engineering*, page 198. 700

Kapitskaia, O., Tomasic, A., and Valduriez, P. (1997). Dealing with discrepancies in wrapper functionality. Research Report RR-3138, INRIA. 319

Karlapalem, K. and Li, Q. (1995). Partitioning schemes for object oriented databases. In *Proc. 5th Int. Workshop on Research Issues on Data Eng.*, pages 42–49. 560

Karlapalem, K., Li, Q., and Vieweg, S. (1996a). Method induced partitioning schemes for object-oriented databases. In *Proc. 16th Int. Conf. on Distributed Computing Systems*, pages 377–384. 564

Karlapalem, K. and Navathe, S. B. (1994). Materialization of redesigned distributed relational databases. Technical Report HKUST-CS94-14, Hong Kong University of Science and Technology, Department of Computer Science. 124

Karlapalem, K., Navathe, S. B., and Ammar, M. (1996b). Optimal redesign policies to support dynamic processing of applications on a distributed relational database system. *Inf. Syst.*, 21(4):353–367. 124

Karlapalem, K., Navathe, S. B., and Morsi, M. A. (1994). Issues in distribution design of object-oriented databases. In Özsu et al. [1994a], pages 148–164. 560

Kashyap, V. and Sheth, A. P. (1996). Semantic and schematic similarities between database objects: A context-based approach. *VLDB J.*, 5(4):276–304. 140, 160

Katz, B. and Lin, J. (2002). Annotating the world wide web using natural language. In *Proc. 2nd Workshop on NLP and XML*, pages 1–8. 681

Katz, H., Chamberlin, D., Draper, D., Fernández, M., Kay, M., Robie, J., Rys, M., Simeon, J., Tivy, J., and Wadler, P. (2004). *XQuery from the Experts: A Guide to the W3C XML Query Language*. Addison Wesley. 719

Kaushik, R., Bohannon, P., Naughton, J. F., and Korth, H. F. (2002). Covering indexing for branching path queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 133–144. 701

Kazerouni, L. and Karlapalem, K. (1997). Stepwise redesign of distributed relational databases. Technical Report HKUST-CS97-12, Hong Kong University of Science and Technology, Department of Computer Science. 124

Keeton, K., Patterson, D., and Hellerstein, J. M. (1998). A case for intelligent disks (idisks). *ACM SIGMOD Rec.*, 27(3):42–52. 499

Keller, A. M. (1982). Update to relational databases through views involving joins. In *Proc. 2nd Int. Conf. on Databases: Improving Usability and Responsiveness*, pages 363–384. 175, 201

Keller, T., Graefe, G., and Maier, D. (1991). Efficient assembly of complex objects. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 148–157. 587, 590, 592

Kementsietsidis, A., Arenas, M., and Miller, R. J. (2003). Managing data mappings in the hyperion project. In *Proc. 19th Int. Conf. on Data Engineering*, pages 732–734. 625

Kemme, B. and Alonso, G. (2000a). Don't be lazy, be consistent: Postgres-R, a new way to implement database replication. In *Proc. 26th Int. Conf. on Very Large Data Bases*, pages 134–143. 482, 548

Kemme, B. and Alonso, G. (2000b). A New Approach to Developing and Implementing Eager Database Replication Protocols. *ACM Trans. Database Syst.*, 25(3):333–379. 482, 548

Kemme, B., Bartoli, A., and O.Babaoglu (2001). Online reconfiguration in replicated databases based on group communication. In *Proc. Int. Conf. on Dependable Systems and Networks*, pages 117–130. 546, 548

Kemme, B., Peris, R. J., and Patino-Martinez, M. (2010). *Database Replication*. Morgan & Claypool. 493

Kemper, A. and Kossmann, D. (1994). Dual-buffering strategies in object bases. In *Proc. 20th Int. Conf. on Very Large Data Bases*, pages 427–438. 570

Kemper, A. and Moerkotte, G. (1990a). Access support in object bases. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 364–374. 587

Kemper, A. and Moerkotte, G. (1990b). Advanced query processing in object bases using access support relations. In *Proc. 16th Int. Conf. on Very Large Data Bases*, pages 290–301. 587

Kemper, A. and Moerkotte, G. (1994). Physical object management. In Kim [1994], pages 175–202. 588, 590, 607

Kermarrec, A.-M., Rowstron, A., Shapiro, M., and Druschel, P. (2001). The icecube approach to the reconciliation of diverging replicas. In *ACM Symp. on Principles of Distributed Computing (PODC)*, pages 210–218. 651

Kermarrec, A.-M. and van Steen, M. (2007). Gossiping in distributed systems. *Operating Systems Rev.*, 41(5):2–7. 617

Kerschberg, L., Ting, P. D., and Yao, S. B. (1982). Query optimization in star computer networks. *ACM Trans. Database Syst.*, 7(4):678–711. 214

Kersten, M. L., Plomp, S., and van den Berg, C. A. (1994). Object storage management in goblin. In Özsu et al. [1994a], pages 100–116. 579

Khoshafian, S. and Copeland, G. (1986). Object identity. In *Proc. Int. Conf. on OOPSLA*, pages 406–416. 553

Khoshafian, S. and Valduriez, P. (1987). Sharing persistence and object-orientation: A database perspective. In *Int. Workshop on Database Programming Languages*, pages 181–205. 251, 292, 510, 553

Kifer, D., Ben-David, S., and Gehrke, J. (2004). Detecting change in data streams. In *Proc. 30th Int. Conf. on Very Large Data Bases*, pages 180–191. 727, 743

Kifer, M., Bernstein, A., and Lewis, P. M. (2006). *Database Systems – An Application-Oriented Approach*. Pearson, 2 edition. 70

Kifer, M., Lausen, G., and Wu, J. (1995). Logical foundations of object-oriented and frame-based languages. *J. ACM*, 42(4):741–843. 607

Kifer, M. and Wu, J. (1993). A logic programming with complex objects. *J. Comp. and System Sci.*, 47(1):77–120. 607

Kim, W. (1984). Highly available systems for database applications. *ACM Comput. Surv.*, 16(1):71–98. 456

Kim, W. (1989). A model of queries for object-oriented databases. In *Proc. 15th Int. Conf. on Very Large Data Bases*, pages 423–432. 587

Kim, W., editor (1994). *Modern Database Management – Object-Oriented and Multidatabase Technologies*. Addison-Wesley/ACM Press. 607, 801

Kim, W., Banerjee, J., Chou, H., Garza, J., and Woelk, D. (1987). Composite objects support in an object-oriented database system. In *Proc. Int. Conf. on OOPSLA*, pages 118–125. 579

Kim, W. and Lochovsky, F., editors (1989). *Object-Oriented Concepts, Databases, and Applications*. Addison Wesley. 607

Kim, W., Reiner, D. S., and Batory, D. S., editors (1985). *Query Processing in Database Systems*. Springer. 220, 807

Kim, W. and Seo, J. (1991). Classifying schematic and data heterogeneity in multidatabase systems. *Comp.*, 24(12):12–18. 160

Kitsuregawa, M. and Ogawa, Y. (1990). Bucket spreading parallel hash: A new, robust, parallel hash join method for data skew in the super database computer. In *Proc. 16th Int. Conf. on Very Large Data Bases*, pages 210–221. 528, 548

Kleinberg, J. (2002). Bursty and hierarchical structure in streams. In *Proc. 8th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, pages 91–101. 727

Kleinberg, J. M. (1999). Authoritative sources in a hyperlinked environment. *J. ACM*, 46(5):604–632. 658, 668

Kleinberg, J. M., Kumar, R., Raghavan, P., Rajagopalan, S., and Tomkins, A. (1999). The Web as a graph: measurements, models, and methods. In *Proc. 5th Annual Int. Conf. Computing and Combinatorics*, pages 1–17. 658

Kling, P., Özsu, M. T., and Daudjee, K. (2010). Distributed XML query processing: Fragmentation, localization and pruning. Technical Report TR-CS-2010-02, University of Waterloo, Cheriton School of Computer Science. 693, 704, 706, 707, 713, 715, 717, 718, 719

Knapp, E. (1987). Deadlock detection in distributed databases. *ACM Comput. Surv.*, 19(4):303–328. 39, 401

Knezevic, P., Wombacher, A., and Risse, T. (2005). Enabling high data availability in a dht. In *Int. Workshop on Grid and P2P Computing Impacts on Large Scale Heterogeneous Distributed Database Systems (GLOBE)*, pages 363–367. 648

Koch, C. (2001). *Data Integration against Multiple Evolving Autonomous Schemata*. Ph.D. thesis, Technical University of Vienna. 133, 134

Koch, C. (2003). Efficient processing of expressive node-selecting queries on XML data in secondary storage: A tree automata-based approach. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 249–260. 700

Kohler, W. H. (1981). A survey of techniques for synchronization and recovery in decentralized computer systems. *ACM Comput. Surv.*, 13(2):149–183. 456

Kollias, J. G. and Hatzopoulos, M. (1981). Criteria to aid in solving the problem of allocating copies of a file in a computer network. *Comp. J.*, 24(1):29–30. 125

Kolodner, E. and Weihl, W. (1993). Atomic incremental garbage collection and recovery for large stable heap. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 177–185. 581

Konopnicki, D. and Shmueli, O. (1995). W3QS: A query system for the World Wide Web. In *Proc. 21th Int. Conf. on Very Large Data Bases*, pages 54–65. 676

Koon, T. M. and Özsu, M. T. (1986). Performance comparison of resilient concurrency control algorithms for distributed databases. In *Proc. 2nd Int. Conf. on Data Engineering*, pages 565–573. 401

Korn, F., Muthukrishnan, S., and Wu, Y. (2006). Modeling skew in data streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 181–192. 727

Korth, H., Levy, E., and Silberschatz, A. (1990). Compensating transactions: A new recovery paradigm. In *Proc. 16th Int. Conf. on Very Large Data Bases*, pages 95–106. 352

Kossmann, D. (2000). The state of the art in distributed query processing. *ACM Comput. Surv.*, 32(4):422–469. 212, 220, 292, 331

Kowalik, J., editor (1985). *Parallel MIMD Computation : the HEP Supercomputer and its applications*. M.I.T. Press. 498

Krämer, J. and Seeger, B. (2005). A temporal foundation for continuous queries over data streams. In *Proc. 11th Int. Conf. on Management of Data (COMAD)*, pages 70–82. 735

Krishnamurthy, R., Boral, H., and Zaniolo, C. (1986). Optimization of non-recursive queries. In *Proc. 11th Int. Conf. on Very Large Data Bases*, pages 128–137. 292

Krishnamurthy, R., Litwin, W., and Kent, W. (1991). Language features for interoperability of databases with schematic discrepancies. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 40–49. 160

Krishnamurthy, S., Franklin, M., Hellerstein, J., and Jacobson, G. (2004). The case for precision sharing. In *Proc. 30th Int. Conf. on Very Large Data Bases*, pages 972–986. 740

Krishnamurthy, S., Wu, C., and Franklin, M. (2006). On-the-fly sharing for streamed aggregation. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 623–634. 741

Krishnaprasad, M., Liu, Z. H., Manikutty, A., Warner, J. W., and Arora, V. (2005). Towards an industrial strength SQL/XML infrastructure. In *Proc. 21st Int. Conf. on Data Engineering*, pages 991–1000. 699

Kshemkalyani, A. and Singhal, M. (1994). On characterization and correctness of distributed deadlocks. *J. Parall. and Distrib. Comput.*, 22(1):44–59. 401

Kubiatowicz, J., Bindel, D., Chen, Y., Czerwinski, S., Eaton, P., Geels, D., Gummadi, R., Rhea, S., Weatherspoon, H., Weimer, W., Wells, C., and Zhao, B. (2000). Oceanstore: an architecture for global-scale persistent storage. In *ACM Int. Conf. on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 190–201. 649, 654

Kumar, A. and Segev, A. (1993). Cost and availability tradeoffs in replicated data concurrency control. *ACM Trans. Database Syst.*, 18(1):102–131. 456, 493

Kumar, R., Raghavan, P., Rajagopalan, S., Sivakumar, D., Tomkins, A., and Upfal, E. (2000). The Web as a graph. In *Proc. 19th ACM SIGACT-SIGMOD-SIGART Symp. on Principles of Database Systems*, pages 1–10. Available from: `http://doi.acm.org/10.1145/335168.335170`. 658, 660

Kumar, R., Raghavan, P., Rajagopalan, S., and Tomkins, A. (1999). Extracting large-scale knowledge bases from the web. In *Proc. 25th Int. Conf. on Very Large Data Bases*, pages 639–650. 660

Kumar, V., editor (1996). *Performance of Concurrency Control Mechanisms in Centralized Database Systems*. Prentice-Hall. 358, 401

Kung, H. T. and Papadimitriou, C. H. (1979). An optimality theory of concurrency control for databases. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 116–125. 350

Kung, H. T. and Robinson, J. T. (1981). On optimistic methods for concurrency control. *ACM Trans. Database Syst.*, 6(2):213–226. 385, 387

Kurose, J. F. and Ross, K. W. (2010). *Computer Networking - A Top-Down Approach Featuring the Internet*. Addison Wesley, 4 edition. 70

Kuss, H. (1982). On totally ordering checkpoint in distributed data bases. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 174–174. 456

Kwok, C. C. T., Etzioni, O., and Weld, D. S. (2001). Scaling question answering to the web. In *Proc. 10th Int. World Wide Web Conf.*, pages 150–161. 681

LaChimia, J. (1984). Query decomposition in a distributed database system using satellite communications. In *Proc. 3rd Seminar on Distributed Data Sharing Systems*, pages 105–118. 214

Lacroix, M. and Pirotte, A. (1977). Domain-oriented relational languages. In *Proc. 3rd Int. Conf. on Very Data Bases*, pages 370–378. 57

Ladin, R. and Liskov, B. (1992). Garbage collection of a distributed heap. In *Proc. 12th Int. Conf. on Distributed Computing Systems*, pages 708–715. 581

Lage, J. P., da Silva, A. S., Golgher, P. B., and Laender, A. H. F. (2002). Collecting hidden weeb pages for data extraction. In *Proc. 4th Int. Workshop on Web Information and Data Management*, pages 69–75. 686

Lakshmanan, L. V. S., Sadri, F., and Subramanian, I. N. (1996). A declarative language for querying and restructuring the Web. In *Proc. 6th Int. Workshop on Research Issues on Data Eng.*, pages 12–21. 676

Lam, K. and Yu, C. T. (1980). An approximation algorithm for a file allocation problem in a hierarchical distributed system. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 125–132. 115

Lam, S. S. and Özsu, M. T. (2002). Querying web data – the WebQA approach. In *Proc. 3rd Int. Conf. on Web Information Systems Eng.*, pages 139–148. 681

Lampson, B. and Sturgis, H. (1976). Crash recovery in distributed data storage system. Technical report, Xerox Palo Alto Research Center, Palo Alto, Calif. 413, 453

Landers, T. and Rosenberg, R. L. (1982). An overview of multibase. In Schneider, H.-J., editor, *Distributed Data Bases*, pages 153–184. North-Holland, Amsterdam. 331

Langville, A. N. and Meyer, C. D. (2006). *Google's PageRank and Beyond*. Princeton University Press. 665

Lanzelotte, R. and Valduriez, P. (1991). Extending the search strategy in a query optimizer. In *Proc. 17th Int. Conf. on Very Large Data Bases*, pages 363–373. 584, 587, 588

Lanzelotte, R., Valduriez, P., and Zaït, M. (1993). On the effectiveness of optimization search strategies for parallel execution spaces. In *Proc. 19th Int. Conf. on Very Large Data Bases*, pages 493–504. 249

Lanzelotte, R., Valduriez, P., Zaït, M., and Ziane, M. (1994). Industrial-strength parallel query optimization: issues and lessons. *Inf. Syst.*, 19(4):311–330. 523, 524, 548

Law, Y.-N., Wang, H., and Zaniolo, C. (2004). Query languages and data models for database sequences and data streams. In *Proc. 30th Int. Conf. on Very Large Data Bases*, pages 492–503. 728

Lawrence, S. and Giles, C. L. (1998). Searching the world wide web. *Science*, 280:98 – 100. 657

Lawrence, S. and Giles, C. L. (1999). Accessibility of information on the web. *Nature*, 400(107 – 109). 657

Lee, M., Freytag, J. C., and Lohman, G. (1988). Implementing an interpreter for functional rules in a query optimizer. In *Proc. 14th Int. Conf. on Very Large Data Bases*, pages 218–229. 586

Lee, S. and Kim, J. (1995). An efficient distributed deadlock detection algorithm. In *Proc. 15th Int. Conf. on Distributed Computing Systems*, pages 169–178. 401

Leland, W., Taqqu, M., Willinger, M., and Wilson, D. (1994). On the self-similar nature of ethernet traffic. *IEEE/ACM Trans. Networking*, 2(1):1–15. 727

Lenoski, D., Laudon, J., Gharachorloo, K., Weber, W. D., Gupta, A., Henessy, J., Horowitz, M., and Lam, M. S. (1992). The stanford dash multiprocessor. *Comp.*, 25(3):63–79. 506, 547

Lenzerini, M. (2002). Data integration: a theoretical perspective. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 233–246. 133

Leon-Garcia, A. and Widjaja, I. (2004). *Communication Networks - Fundamental Concepts and Key Architectures*. McGraw-Hill, 2 edition. 70

Leung, J. Y. and Lai, E. K. (1979). On minimum cost recovery from system deadlock. *IEEE Trans. Comput.*, 28(9):671–677. 391

Levin, K. D. and Morgan, H. L. (1975). Optimizing distributed data bases: A framework for research. In *Proc. National Computer Conf*, pages 473–478. 38, 71, 125

Levy, A. Y., Mendelzon, A. O., Sagiv, Y., and Srivastava, D. (1995). Answering queries using views. In *Proc. ACM SIGACT-SIGMOD Symp. on Principles of Database Systems*, pages 95–104. 304, 331

Levy, A. Y., Rajaraman, A., and Ordille, J. J. (1996a). Querying heterogeneous information sources using source descriptions. In *Proc. 22th Int. Conf. on Very Large Data Bases*, pages 251–262. 160

Levy, A. Y., Rajaraman, A., and Ordille, J. J. (1996b). Querying heterogeneous information sources using source descriptions. In *Proc. 22th Int. Conf. on Very Large Data Bases*, pages 251–262. 305, 331

Levy, A. Y., Rajaraman, A., and Ordille, J. J. (1996c). The world wide web as a collection of views: Query processing in the information manifold. In *Proc. Workshop on Materialized Views: Techniques and Applications*, pages 43–55. 160

Li, F., Chang, C., Kollios, G., and Bestavros, A. (2006). Characterizing and exploiting reference locality in data stream applications. In *Proc. 22nd Int. Conf. on Data Engineering*, page 81. 740

Li, V. O. K. (1987). Performance models of timestamp-ordering concurrency control algorithms in distributed databases. *IEEE Trans. Comput.*, C-36(9):1041–1051. 401

Li, W.-S. and Clifton, C. (2000). Semint: A tool for identifying attribute correspondences in heterogeneous databases using neural networks. *Data & Knowl. Eng.*, 33(1):49–84. 145

Li, W.-S., Clifton, C., and Liu, S.-Y. (2000). Database integration using neural networks: Implementation and experiences. *Knowl. and Information Syst.*, 2(1):73–96. 145

Liang, D. and Tripathi, S. K. (1996). Performance analysis of long-lived transaction processing systems with rollbacks and aborts. *IEEE Trans. Knowl. and Data Eng.*, 8(5):802–815. 401

Lim, H.-S., Lee, J.-G., Lee, M.-J., Whang, K.-Y., and Song, I.-Y. (2006). Continuous query processing in data streams using duality of data and queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 313–324. 741

Lim, L., Wang, M., Padmanabhan, S., Vitter, J. S., and Agarwal, R. (2003). Dynamic maintenance of web indexes using landmarks. In *Proc. 12th Int. World Wide Web Conf.*, pages 102–111. 667

Lima, A., Mattoso, M., and Valduriez, P. (2004a). Olap query processing in a database cluster. In *Proc. 10th Int. Euro-Par Conf.*, pages 355–362. 543, 548

Lima, A. A. B., Mattoso, M., and Valduriez, P. (2004b). Adaptive virtual partitioning for olap query processing in a database cluster. In *Proc. Brazilian Symposium on Databases*, pages 92–105. 544, 548

Lin, W. K. (1981). Performance evaluation of two concurrency control mechanisms in a distributed database system. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 84–92. 401

Lin, W. K. and Nolte, J. (1982). Performance of two phase locking. In *Proc. 6th Berkeley Workshop on Distributed Data Management and Computer Networks*, pages 131–160. 401

Lin, W. K. and Nolte, J. (1983). Basic timestamp, multiple version timestamp, and two-phase locking. In *Proc. 9th Int. Conf. on Very Data Bases*, pages 109–119. 401

Lin, X., Lu, H., Xu, J., and Yu, J. X. (2004). Continuously maintaining quantile summaries of the most recent N elements over a data stream. In *Proc. 20th Int. Conf. on Data Engineering*, pages 362–373. 727, 737

Lin, Y., Kemme, B., Patiño-Martínez, M., and Jiménez-Peris, R. (2005). Middleware based data replication providing snapshot isolation. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 419–430. 464

Lindsay, B. (1979). Notes on distributed databases. Technical Report RJ 2517, IBM San Jose Research Laboratory, San Jose, Calif. 426

Liskov, B., Adya, A., Castro, M., Day, M., Ghemawat, S., Gruber, R., Maheshwari, U., Myers, A., and Shrira, L. (1996). Safe and efficient sharing of persistent objects in thor. In *ACM SIGMOD Int. Conf. on Management of Data*, pages 318–329. 568, 569

Liskov, B., Day, M., and Shirira, L. (1994). Distributed object management in thor. In Özsu et al. [1994a], pages 79–91. 577

Litwin, W. (1988). From database systems to multidatabase systems: Why and how. In *Proc. British National Conference on Databases*, pages 161–188, Cambridge. Cambridge University Press. 40

Litwin, W., Neimat, M.-A., and Schneider, D. A. (1993). LH* – linear hashing for distributed files. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 327–336. 618

Liu, B., Zhu, Y., and Rundensteiner, E. (2006). Run-time operator state spilling for memory intensive long running queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 347–358. 740

Liu, L., Pu, C., Barga, R., and Zhou, T. (1996). Differential evaluation of continual queries. In *Proc. IEEE Int. Conf. Dist. Comp. Syst*, pages 458–465. 6

Liu, L., Pu, C., and Tang, W. (1999). Continual queries for internet-scale event-driven information delivery. *IEEE Trans. Knowl. and Data Eng.*, 11(4):610–628. 736

Liu, Z. H., Chandrasekar, S., Baby, T., and Chang, H. J. (2008). Towards a physical XML independent XQuery/sql/xml engine. *Proc. VLDB*, 1(2):1356–1367. 698

Livny, M., Khoshafian, S., and Boral, H. (1987). Multi-disk management. In *Proc. ACM SIGMETRICS Conf. on Measurement and Modeling of Computer Systems*, pages 69–77. 508, 510, 548

Lohman, G. and Mackert, L. F. (1986). R* optimizer validation and performance evaluation for distributed queries. In *Proc. 11th Int. Conf. on Very Large Data Bases*, pages 149–159. 281, 293

Lohman, G., Mohan, C., Haas, L., Daniels, D., Lindsay, B., Selinger, P., and Wilms, P. (1985). Query processing in r*. In Kim et al. [1985], pages 31–47. 250, 277

Longbottom, R. (1980). *Computer System Reliability*. John Wiley & Sons. 410, 455

Lu, H. and Carey, M. J. (1985). Some experimental results on distributed join algorithms in a local network. In *Proc. 10th Int. Conf. on Very Large Data Bases*, pages 292–304. 273

Lu, H., Ooi, B., and Goh, C. (1992). On global multidatabase query optimization. *ACM SIGMOD Rec.*, 21(4):6–11. 307, 331

Lu, H., Ooi, B., and Goh, C. (1993). Multidatabase query optimization: Issues and solutions. In *Proc. 3rd Int. Workshop on Res. Issues in Data Eng*, pages 137–143. 298, 331

Lu, H., Shan, M.-C., and Tan, K.-L. (1991). Optimization of multi-way join queries for parallel execution. In *Proc. 17th Int. Conf. on Very Large Data Bases*, pages 549–560. 530

Lunt, T. F., Denning, D. E., Schell, R. R., Heckman, M., and Shockley, W. R. (1990). The SeaView security model. *IEEE Trans. Softw. Eng.*, 16(6):593–607. 184

Lunt, T. F. and Fernández, E. B. (1990). Database security. *ACM SIGMOD Rec.*, 19(4):90–97. 181, 201, 202

Lv, Q., Cao, P., Cohen, E., Li, K., and Shenker, S. (2002). Search and replication in unstructured peer-to-peer networks. In *Proc. 16th Annual Int. Conf. on Supercmputing*, pages 84–95. 617

Lynch, N. (1983a). Concurrency control for resilient nested transactions. In *Proc. 2nd ACM SIGACT–SIGMOD Symp. on Principles of Database Systems*, pages 166–181. 401

Lynch, N. (1983b). Multilevel atomicity: A new correctness criterion for database concurrency control. *ACM Trans. Database Syst.*, 8(4):484–502. 395, 401

Lynch, N. and Merritt, M. (1986). Introduction to the theory of nested transactions. Technical Report MIT/LCS/TR-367, Massachusetts Institute of Technology, Cambridge, Mass. 401

Lynch, N., Merritt, M., Weihl, W. E., and Fekete, A. (1993). *Atomic Transactions in Concurrent Distributed Systems*. Morgan Kaufmann. 359, 401

Ma, L., Viglas, S., Li, M., and Li, Q. (2005). Stream operators for querying data streams. In *Proc. 6th Int. Conf. on Web-Age Information Management:*, pages 404–415. 727

Mackert, L. F. and Lohman, G. (1986). R* optimizer validation and performance evaluation for local queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 84–95. 264, 281, 291

Madden, S. and Franklin, M. J. (2002). Fjording the stream: An architecture for queries over streaming sensor data. In *Proc. 18th Int. Conf. on Data Engineering*, pages 555–566. 734

Madden, S., Shah, M., Hellerstein, J., and Raman, V. (2002a). Continuously adaptive continuous queries over streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 49–60. 734, 741

Madden, S., Shah, M. A., Hellerstein, J. M., and Raman, V. (2002b). Continuously adaptive continuous queries over streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 49–60. 320

Madhavan, J., Bernstein, P. A., and Rahm, E. (2001). Generic schema matching with cupid. In *Proc. 27th Int. Conf. on Very Large Data Bases*, pages 49–58. 134, 144, 160

Maheshwari, U. and Liskov, B. (1994). Fault-tolerant distributed garbage collection in a client-server object-oriented database. In *Proc. 3rd Int. Conf. on Parallel and Distributed Information Systems*, pages 239–248. 581

Mahmoud, . A. and Riordon, J. S. (1976). Optimal allocation of resources in distributed information networks. *ACM Trans. Database Syst.*, 1(1):66–78. 125

Maier, D. (1986). A logic for objects. Technical Report CS/E-86-012, Oregon Graduate Center. 607

Maier, D. (1989). Why isn't there an object-oriented data model? Technical Report CS/E 89-002, Oregon Graduate Center, Portland, Oregon. 553

Maier, D., Graefe, G., Shapiro, L., Daniels, S., Keller, T., and Vance, B. (1994). Issues in distributed object assembly. In Özsu et al. [1994a], pages 165–181. 592

Maier, D. and Stein, J. (1986). Indexing in an object-oriented dbms. In *Proc. Int. Workshop on Object-Oriented Database Systems*, pages 171–182. 587, 588, 589, 590

Makki, K. and Pissinou, N. (1995). Detection and resolution algorithm for deadlocks in distributed database systems. In *Proc. ACM Int. Conf. on Information and Knowledge Management*, pages 411–416. 401

Malkhi, D., Noar, M., and Ratajczak, D. (2002). Viceroy: A scalable and dynamic emulation of the butterfly. In *Proc. ACM SIGACT-SIGOPS 21st Symp. on the Principles of Distributed Computing*, pages 183–192. 621

Manber, U. and Myers, G. (1990). Suffix arrays: a new method for on-line string searches. In *Proc. 1st Annual ACM-SIAM Symp. on Discrete Algorithms*, pages 319–327. 667

Manolescu, I., Florescu, D., and Kossmann, D. (2001). Answering XML queries on heterogeneous data sources. In *Proc. 27th Int. Conf. on Very Large Data Bases*, pages 241–250. 160

Martin, B. and Pedersen, C. H. (1994). Long-lived concurrent activities. In Özsu et al. [1994a], pages 188–211. 593

Martínez, J. M., editor. MPEG-7 overview (2004). Available from: `http://www.chiariglione.org/mpeg/standards/mpeg-7/mpeg-7.htm` [Last retrieved: December 2009]. 690

Martins, V., Akbarinia, R., Pacitti, E., and Valduriez, P. (2006a). Reconciliation in the appa p2p system. In *IEEE Int. Conf. on Parallel and Distributed Systems (ICPADS)*, pages 401–410. 651, 654

Martins, V. and Pacitti, E. (2006). Dynamic and distributed reconciliation in p2p-dht networks. In *uropean Conf. on Parallel Computing (Euro-Par)*, pages 337–349. 651, 654

Martins, V., Pacitti, E., Dick, M. E., and Jimenez-Peris, R. (2008). Scalable and topology-aware reconciliation on p2p networks. *Distrib. Parall. Databases*, 24(1–3):1–43. 651

Martins, V., Pacitti, E., and Valduriez, P. (2006b). Survey of data replication in p2p systems. Technical Report 6083, INRIA, Rennes, France. 654

Maymounkov, P. and Mazières, D. (2002). Kademlia: A peer-to-peer information system based on the XOR metric. In *Proc. 1st Int. Workshop Peer-to-Peer Systems*, Lecture Notes in Computer Science 2429, pages 53–65. 621

McBrien, P. and Poulovassilis, A. (2003). Defining peer-to-peer data integration using both as view rules. In *Proc. 1st Int. Workshop on Databases, Information Systems and Peer-to-Peer Computing*, pages 91–107. 627

McCallum, A., Nigam, K., Rennie, J., and Seymore, K. (1999). A machine learning approach to building domain-specific search engines. In *Proc. 16th Int. Joint Conf. on AI*. 666

McCann, R., AlShebli, B., Le, Q., Nguyen, H., Vu, L., and Doan, A. (2005). Mapping maintenance for data integration systems. In *Proc. 31st Int. Conf. on Very Large Data Bases*, pages 1018–1029. 156

McConnel, S. and Siewiorek, D. P. (1982). Evaluation criteria. In Siewiorek and Swarz [1982], pages 201–302. 409

McCormick, W. T., Schweitzer, P. J., and White, T. W. (1972). Problem decomposition and data reorganization by a clustering technique. *Oper. Res.*, 20(5):993–1009. 102

Medina-Mora, R., Wong, H., and Flores, P. (1993). Action workflow as the enterprise integration technology. *Q. Bull. IEEE TC on Data Eng.*, 16(2):49–52. 354

Mehta, M. and DeWitt, D. (1995). Managing intra-operator parallelism in parallel database systems. In *Proc. 21th Int. Conf. on Very Large Data Bases*. 529, 548

Melnik, S., Garcia-Molina, H., and Rahm, E. (2002). Similarity flooding: A versatile graph matching algorithm and its application to schema matching. In *Proc. 18th Int. Conf. on Data Engineering*, pages 117–128. 134, 145, 148, 160

Melnik, S., Raghavan, S., Yang, B., and Garcia-Molina, H. (2001). Building a distributed full-text index for the web. In *Proc. 10th Int. World Wide Web Conf.*, pages 396–406. Available from: `citeseer.ist.psu.edu/article/melnik01building.html`. 668

Melton, J. (2002). *Advanced SQL: 1999 - Understanding Object-Relational and Other Advanced Features*. Morgan Kaufmann. 553

Melton, J., Michels, J.-E., Josifovski, V., Kulkarni, K., Schwartz, P., and Zeidenstein, K. (2001). Sql and management of external data. *ACM SIGMOD Rec.*, 30(1):70–77. 314, 328

Menasce, D. A. and Muntz, R. R. (1979). Locking and deadlock detection in distributed databases. *IEEE Trans. Softw. Eng.*, SE-5(3):195–202. 392

Menasce, D. A. and Nakanishi, T. (1982a). Optimistic versus pessimistic concurrency control mechanisms in database management systems. *Inf. Syst.*, 7(1):13–27. 401

Menasce, D. A. and Nakanishi, T. (1982b). Performance evaluation of a two-phase commit based protocol for ddbs. In *Proc. First ACM SIGACT–SIGMOD Symp. on Principles of Database Systems*, pages 247–255. 401

Mendelzon, A. O., Mihaila, G. A., and Milo, T. (1997). Querying the World Wide Web. *Int. J. Digit. Libr.*, 1(1):54–67. 676, 677

Meng, W., Yu, C., Kim, W., Wang, G., Phan, T., and Dao, S. (1993). Construction of relational front-end for object-oriented database systems. In *Proc. 9th Int. Conf. on Data Engineering*, pages 476–483. 331

Merrett, T. H. and Rallis, N. (1985). An analytic evaluation of concurrency control algorithms. In *Proc. CIPS (Canadian Information Processing Society) Congress '85*, pages 435–439. 401

Milán-Franco, J. M., Jiménez-Peris, R., Patiño-Martínez, M., and Kemme, B. (2004). Adaptive middleware for data replication. In *Proc. ACM/IFIP/USENIX Int. Middleware Conf.*, pages 175–194. 542, 548

Miller, G. A. (1995). WordNet: A lexical database for English. *Commun. ACM*, 38(11):39–45. 142

Miller, R. J., Haas, L. M., and Hernández, M. A. (2000). Schema mapping as query discovery. In *Proc. 26th Int. Conf. on Very Large Data Bases*, pages 77–88. 150

Miller, R. J., Hernández, M. A., Haas, L. M., Yan, L., Ho, C. T. H., Fagin, R., and Popa, L. (2001). The Clio project: Managing heterogeneity. *ACM SIGMOD Rec.*, 31(1):78–83. 152

Milo, T. and Suciu, D. (1999). Index structures for path expressions. In *Proc. 7th Int. Conf. on Database Theory*, pages 277–295. 701

Milo, T. and Zohar, S. (1998). Using schema matching to simplify heterogeneous data translation. In *Proc. 24th Int. Conf. on Very Large Data Bases*, pages 122–133. 134, 160

Minoura, T. and Wiederhold, G. (1982). Resilient extended true-copy token scheme for a distributed database system. *IEEE Trans. Softw. Eng.*, SE-8(3):173–189. 456, 493

Mitchell, G., Dayal, U., and Zdonik, S. (1993). Control of an extensible query optimizer: A planning-based approach. In *Proc. 19th Int. Conf. on Very Large Data Bases*, pages 517–528. 584

Mitchell, T. (1997). *Machine Learning*. McGraw-Hill. 666

Mohan, C. (1979). Data base design in the distributed environment. Working Paper WP-7902, Department of Computer Sciences, University of Texas at Austin. 125

Mohan, C. and Lindsay, B. (1983). Efficient commit protocols for the tree of processes model of distributed transactions. In *Proc. ACM SIGACT-SIGOPS 2nd Symp. on the Principles of Distributed Computing*, pages 76–88. 434, 456

Mohan, C., Lindsay, B., and Obermarck, R. (1986). Transaction management in the r* distributed database management system. *ACM Trans. Database Syst.*, 11(4):378–396. 377, 393, 434

Mohan, C. and Yeh, R. T. (1978). *Distributed Data Base Systems: A Framework for Data Base Design. In Distributed Data Bases, Infotech State-of-the-Art Report*. Infotech. 39

Morgan, H. L. and Levin, K. D. (1977). Optimal program and data location in computer networks. *Commun. ACM*, 20(5):315–322. 125

Moss, E. (1985). *Nested Transactions*. M.I.T. Press. 351, 352, 396, 401

Motwani, R., Widom, J., Arasu, A., Babcock, B., Babu, S., Datar, M., Manku, G., Olston, C., Rosenstein, J., and Varma, R. (2003). Query processing, approximation, and resource management in a data stream management system. In *Proc. 1st Biennial Conf. on Innovative Data Systems Research*, pages 245–256. 732

Muro, S., Ibaraki, T., Miyajima, H., and Hasegawa, T. (1983). File redundancy issues in distributed database systems. In *Proc. 9th Int. Conf. on Very Data Bases*, pages 275–277. 124

Muro, S., Ibaraki, T., Miyajima, H., and Hasegawa, T. (1985). Evaluation of file redundancy in distributed database systems. *IEEE Trans. Softw. Eng.*, SE-11(2):199–205. 124

Muth, P., Rakow, T., Weikum, G., Brössler, P., and Hasse, C. (1993). Semantic concurrency control in object-oriented database systems. In *Proc. 9th Int. Conf. on Data Engineering*, pages 233–242. 604, 605

Myers, G. J. (1976). *Software Reliability: Principles and Practices*. John Wiley & Sons. 455

Naacke, H., Tomasic, A., and Valduriez, P. (1999). Validating mediator cost models with DISCO. *Networking and Information Systems Journal*, 2(5):639–663. 307, 310, 331

Najork, M. and Wiener, J. L. (2001). Breadth-first crawling yields high-quality pages. In *Proc. 10th Int. World Wide Web Conf.*, pages 114–118. 665

Naumann, F., Ho, C.-T., Tian, X., Haas, L. M., and Megiddo, N. (2002). Attribute classification using feature analysis. In *Proc. 18th Int. Conf. on Data Engineering*, page 271. 146

Navathe, S. B., Ceri, S., Wiederhold, G., and Dou, J. (1984). Vertical partitioning of algorithms for database design. *ACM Trans. Database Syst.*, 9(4):680–710. 98, 99, 102, 109, 125

NBS (1977). Data encryption standard. Technical Report 46, U. S. Department of Commerce/National Bureau of Standards, Federal Information Processing Standards Publication. 180

Nejdl, W., Siberski, W., and Sintek, M. (2003). Design issues and challenges for rdf- and schema-based peer-to-peer systems. *ACM SIGMOD Rec.*, 32(3):41–46. 624, 628

Nepal, S. and Ramakrishna, M. (1999). Query processing issues in image (multimedia) databases. In *Proc. 15th Int. Conf. on Data Engineering*, pages 22–29. 629, 654

Newton, G. (1979). Deadlock prevention, detection and resolution: An annotated bibliography. *Operating Systems Rev.*, 13(2):33–44. 401

Ng, P. (1988). A commit protocol for checkpointing transactions. In *Proc. 7th. Symp. on Reliable Distributed Systems*, pages 22–31. 456

Niamir, B. (1978). Attribute partitioning in a self–adaptive relational database system. Technical Report 192, Laboratory for Computer Science, Massachusetts Institute of Technology, Cambridge, Mass. 98, 125

Nicola, M. and der Linden, B. V. (2005). Native XML support in db2 universal database. In *Proc. 31st Int. Conf. on Very Large Data Bases*, pages 1164–1174. 699

Nicolas, J. M. (1982). Logic for improving integrity checking in relational data bases. *Acta Informatica*, 18:227–253. 192, 202

Nodine, M. and Zdonik, S. (1990). Cooperative transaction hierarchies: A transaction model to support design applications. In *Proc. 16th Int. Conf. on Very Large Data Bases*, pages 83–94. 354

OASIS UDDI. Universal description discovery & integration (UDDI) (2002). Available from: `http://uddi.xml.org/` [Last retrieved: December 2009]. 690

Obermarck, R. (1982). Deadlock detection for all resource classes. *ACM Trans. Database Syst.*, 7(2):187–208. 39, 393, 401

Omiecinski, E. (1991). Performance analysis of a load balancing hash-join algorithm for a shared-memory multiprocessor. In *Proc. 17th Int. Conf. on Very Large Data Bases*, pages 375–385. 528, 548

Ooi, B., Shu, Y., and Tan, K.-L. (2003a). Relational data sharing in peer-based data management systems. *ACM SIGMOD Rec.*, 32(3):59–64. 627

Ooi, B. C., Shu, Y., and Tan, K.-L. (2003b). Db-enabled peers for managing distributed data. In *Proc. 5th Asian-Pacific Web Conference*, pages 10–21. 612

Ordonez, C. (2003). Clustering binary data streams with k-means. In *Proc. ACM SIGMOD Workshop on Research Issues in Data Mining and Knowledge Discovery*. 743

Orenstein, J., Haradvala, S., Margulies, B., and Sakahara, D. (1992). Query processing in the objectstore database system. In *ACM SIGMOD Int. Conf. on Management of Data*, pages 403–412. 586

Orfali, R., Harkey, D., and Edwards, J. (1996). *The Essential Distributed Objects Survival Guide*. John Wiley & Sons. 607

Osborn, S. L. and Heaven, T. E. (1986). The design of a relational database system with abstract data types for domains. *ACM Trans. Database Syst.*, 11(3):357–373. 557

Osterhaug, A. (1989). *Guide to Parallel Programming on Sequent Computer Systems*. Prentice-Hall. 498

O'Toole, J., Nettles, S., and Gifford, D. (1993). Concurrent compacting garbage collection of a persistent heap. In *Proc. 14th ACM Symp. Operating Syst. Principles*, pages 161–174. 581

Ou, Z., Yu, G., Yu, Y., Wu, S., Yang, X., and Deng, Q. (2005). Tick scheduling: A deadline based optimal task scheduling approach for real-time data stream systems. In *Proc. 6th Int. Conf. on Web-Age Information Management:*, pages 725–730. 735

Ouksel, A. M. and Sheth, A. P. (1999). Semantic interoperability in global information systems: A brief introduction to the research area and the special section. *ACM SIGMOD Rec.*, 28(1):5–12. 160

Özsoyoglu, Z. M. and Zhou, N. (1987). Distributed query processing in broadcasting local area networks. In *Proc. 20th Hawaii Int. Conf. on System Sciences*, pages 419–429. 214, 215

Özsu, M. and Barker, K. (1990). Architectural classification and transaction execution models of multidatabase systems. In *Proc. Int. Conf. on Computing and Information*, pages 275–279. 40

Özsu, M., Dayal, U., and Valduriez, P., editors (1994a). *Distributed Object Management*. Morgan Kaufmann, San Mateo, Calif. 607, 784, 785, 787, 789, 793, 800, 801, 807, 809, 814

Özsu, M., Peters, R., Szafron, D., Irani, B., Munoz, A., and Lipka, A. (1995a). Tigukat: A uniform behavioral objectbase management system. *VLDB J.*, 4:445–492. 555, 606

Özsu, M. T. (1985a). Modeling and analysis of distributed concurrency control algorithms using an extended petri net formalism. *IEEE Trans. Softw. Eng.*, SE-11(10):1225–1240. 401

Özsu, M. T. (1985b). Performance comparison of distributed vs centralized locking algorithms in distributed database systems. In *Proc. 5th Int. Conf. on Distributed Computing Systems*, pages 254–261. 401

Özsu, M. T. (1994). Transaction models and transaction management in OODBMSs. In Dogac et al. [1994], pages 275–279. 359, 607

Özsu, M. T. and Blakeley, J. (1994). Query processing in object-oriented database systems. In Kim, W., editor, *Modern Database Management – Object-Oriented and Multidatabase Technologies*, pages 146–174. Addison-Wesley/ACM Press. 582, 607

Özsu, M. T., Dayal, U., and Valduriez, P. (1994b). An introduction to distributed object management. In Özsu et al. [1994a], pages 1–24. 551

Özsu, M. T., Munoz, A., and Szafron, D. (1995b). An extensible query optimizer for an objectbase management system. In *Proc. 4th Int. Conf. on Information and Knowledge Management*, pages 188–196. 584

Özsu, M. T. and Valduriez, P. (1991). Distributed database systems: Where are we now? *Comp.*, 24(8):68–78. 38

Özsu, M. T. and Valduriez, P. (1994). Distributed data management: Unsolved problems and new issues. In Casavant, T. and Singhal, M., editors, *Readings in Distributed Computing Systems*, pages 512–544. IEEE/CS Press. 38

Özsu, M. T. and Valduriez, P. (1997). Distributed and parallel database systems. In Tucker, A., editor, *Handbook of Computer Science and Engineering*, pages 1093–1111. CRC Press. 38

Özsu, M. T., Voruganti, K., and Unrau, R. (1998). An asynchronous avoidance-based cache consistency algorithm for client caching dbmss. In *Proc. 24th Int. Conf. on Very Large Data Bases*, pages 440–451. 573

Pacitti, E., Coulon, C., Valduriez, P., and Özsu, M. T. (2006). Preventive replication in a database cluster. *Distrib. Parall. Databases*, 18(3):223–251. 537, 539, 540, 548

Pacitti, E., Minet, P., and Simon, E. (1999). Fast algorithms for maintaining replica consistency in lazy master replicated databases. In *Proc. 25th Int. Conf. on Very Large Data Bases*, pages 126–137. 463, 482, 484, 537

Pacitti, E., Özsu, M. T., and Coulon, C. (2003). Preventive multi-master replication in a cluster of autonomous databases. In *Proc. 9th Int. Euro-Par Conf.*, pages 318–327. 537, 548

Pacitti, E. and Simon, E. (2000). Update propagation strategies to improve freshness in lazy master replicated databases. *VLDB J.*, 8(3-4):305–318. 462, 493, 537

Pacitti, E., Simon, E., and de Melo, R. (1998). Improving data freshness in lazy master schemes. In *Proc. 18th Int. Conf. on Distributed Computing Systems*, pages 164–171. 463, 493

Pacitti, E., Valduriez, P., and Mattoso, M. (2007a). Grid data management: open problems and new issues. *Journal of Grid Computing*, 5(3):273–281. 654

Pacitti, E., Valduriez, P., and Mattoso, M. (2007b). Grid data management: Open problems and new issues. *J. Grid Comp.*, 5(3):273–281. 750, 763

Page, L., Brin, S., Motwani, R., and Winograd, T. (1998). The pagerank citation ranking: Bringing order to the web. Technical report, Stanford University. 665

Page, T. W. and Popek, G. J. (1985). Distributed data management in local area networks. In *Proc. ACM SIGACT–SIGMOD Symp. on Principles of Database Systems*, pages 135–142. 210, 250

Pal, S., Cseri, I., Seeliger, O., Rys, M., Schaller, G., Yu, W., Tomic, D., Baras, A., Berg, B., Churin, D., and Kogan, E. (2005). Xquery implementation in a relational database system. In *Proc. 31st Int. Conf. on Very Large Data Bases*, pages 1175–1186. 699

Palma, W., Akbarinia, R., Pacitti, E., and Valduriez, P. (2009). Dhtjoin: processing continuous join queries using dht networks. *Distrib. Parall. Databases*, 26(2–3):291–317. 732

Palopoli, L., Saccà, D., Terracina, G., and Ursino, D. (1999). A unified graph-based framework for deriving nominal interscheme properties, type conflicts and object cluster similarities. In *Proc. Int. Conf. on Cooperative Information Systems*, pages 34–45. 134, 142, 160

Palopoli, L., Saccà, D., Terracina, G., and Ursino, D. (2003a). Uniform techniques for deriving similarities of objects and subschemes in heterogeneous databases. *IEEE Trans. Knowl. and Data Eng.*, 15(2):271–294. 145, 160

Palopoli, L., Saccà, D., and Ursino, D. (1998). Semi-automatic semantic discovery of properties from database schemas. In *Proc. Int. Conf. on Database Eng. and Applications*, pages 244–253. 134, 145, 160

Palopoli, L., Terracina, G., and Ursino, D. (2003b). Experiences using DIKE, a system for supporting cooperative information system and data warehouse design. *Inf. Syst.*, 28:835–865. 134, 160

Palpanas, T., Vlachos, M., Keogh, E., Gunopulos, D., and Truppel, W. (2004). Online amnesic approximation of streaming time series. In *Proc. 20th Int. Conf. on Data Engineering*, pages 338–349. 726

Pandey, S., Ramamritham, K., and Chakrabarti, S. (2003). Monitoring the dynamic web to respond to continuous queries. In *Proc. 12th Int. World Wide Web Conf.* 6

Papadimitriou, C. H. (1979). Serializability of concurrent database updates. *J. ACM*, 26(4):631–653. 350

Papadimitriou, C. H. (1986). *The Theory of Concurrency Control*. Computer Science Press. 401

Papakonstantinou, Y., Garcia-Molina, H., and Widom, J. (1995). Object exchange across heterogeneous information sources. In *Proc. 11th Int. Conf. on Data Engineering*, pages 251–260. 671, 673

Pape, C. L., Gançarski, S., and Valduriez, P. (2004). Refresco: Improving query performance through freshness control in a database cluster. In *Proc. Confederated Int. Conf. DOA, CoopIS and ODBASE*, Lecture Notes in Computer Science 3290, pages 174–193. 493, 540, 548

Paris, J. F. (1986). Voting with witnesses: A consistency scheme for replicated files. In *Proc. 6th Int. Conf. on Distributed Computing Systems*, pages 606–612. 493

Park, Y., Scheuermann, P., and Tang, H. (1995). A distributed deadlock detection and resolution algorithm based on a hybrid wait-for graph and probe generation scheme. In *Proc. ACM Int. Conf. Information and Knowledge Management*, pages 378–86. 401

Passerini, A., Frasconi, P., and Soda, G. (2001). Evaluation methods for focused crawling. In *Proc. 7th Congress of the Italian Association for Artificial Intelligence*, pages 33–39. 666

Patiño-Martínez, M., Jiménez-Peris, R., Kemme, B., and Alonso, G. (2005). MIDDLE-R: Consistent database replication at the middleware level. *ACM Trans. Comp. Syst.*, 23(4):375–423. 491

Patiño-Martínez, M., Jiménez-Peris, R., Kemme, B., and Alonso, G. (2000). Scalable replication in database clusters. In *Proc. 14th Int. Symp. on Distributed Computing*, pages 315–329. 482, 489, 548

Pavlo, A., Paulson, E., Rasin, A., Abadi, D. J., DeWitt, D. J., Madden, S., and Stonebraker, M. (2009). A comparison of approaches to large-scale data analysis. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 165–178. 760

Paxson, V. and Floyd, S. (1995). Wide-area traffic: The failure of poisson modeling. *IEEE/ACM Trans. Networking*, 3(3):226–244. 727

Pease, M., Shostak, R., and Lamport, L. (1980). Reaching agreement in the presence of faults. *J. ACM*, 27(2):228–234. 456

Pedone, F. and Schiper, A. (1998). Optimistic atomic broadcast. In *Proc. 12th Int. Symp. on Distributed Computing*, pages 318–332. 539

Perez-Sorrosal, F., Vuckovic, J., Patiño-Martínez, M., and Jiménez-Peris, R. (2006). Highly available long running transactions and activities for J2EE. In *Proc. 26th Int. Conf. on Distributed Computing Systems*, page 2. 546, 548

Peters, R. J., Lipka, A., Özsu, M. T., and Szafron, D. (1993). An extensible query model and its languages for a uniform behavioral object management system. In *Proc. 2nd International Conference on Information and Knowledge Management*, pages 403–412. 584

Piatetsky-Shapiro, G. and Connell, C. (1984). Accurate estimation of the number of tuples satisfying a condition. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 256–276. 252

Pinedo, M. (2001). *Scheduing: Theory, Algorithms and Systems*. Integre Technical Publishing, 2 edition. 537

Pirahesh, H., Mohan, C., Cheng, J. M., Liu, T. S., and Selinger, P. G. (1990). Parallelism in rdbms : Architectural issues and design. In *Proc. 2nd Int. Symp. on Databases in Distributed and Parallel Systems*, pages 4–29. 532, 533, 548

Plainfossé, D. and Shapiro, M. (1995). A survey of distributed garbage collection techniques. In *Proc. Int. Workshop on Memory Management*, pages 211–249. 581

Plattner, C. and Alonso, G. (2004). Ganymed: Scalable replication for transactional web applications. In *Proc. ACM/IFIP/USENIX Int. Middleware Conf.*, pages 155–174. 464

Plaxton, C., Rajaraman, R., and Richa, A. (1997). Accessing nearby copies of replicated objects in a distributed environment. In *ACM Symp. on Parallel Algorithms and Architectures (SPAA)*, pages 311–320. 646

Polyzotis, N. and Garofalakis, M. N. (2002). Statistical synopses for graph-structured XML databases. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 358–369. 701

Polyzotis, N., Garofalakis, M. N., and Ioannidis, Y. E. (2004). Approximate XML query answers. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 263–274. 701

Polyzotis, N., Skiadopoulos, S., Vassiliadis, P., Simitsis, A., and Frantzell, N.-E. (2008). Meshing streaming updates with persistent data in an active data warehouse. *IEEE Trans. Knowl. and Data Eng.*, 20(7):976–991. 761

Poosala, V., Ioannidis, Y., Haas, P., and Shekita, E. (1996). Improved histograms for selectivity estimation of range predicates. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 294–305. 256

Popa, L., Velegrakis, Y., Miller, R. J., Hernandez, M. A., and Fagin, R. (2002). Translating web data. In *Proc. 28th Int. Conf. on Very Large Data Bases*. 155

Porto, F., Laber, E. S., and Valduriez, P. (2003). Cherry picking: A semantic query processing strategy for the evaluation of expensive predicates. In *Proc. Brazilian Symposium on Databases*, pages 356–370. 320, 326, 331

Potier, D. and LeBlanc, P. (1980). Analysis of locking policies in database management systems. *Commun. ACM*, 23(10):584–593. 401

Pottinger, R. and Levy, A. Y. (2000). A scalable algorithm for answering queries using views. In *Proc. 26th Int. Conf. on Very Large Data Bases*, pages 484–495. 305, 331

Pradhan, D. K., editor (1986). *Fault-Tolerant Computing: Theory and Techniques*, volume 2. Prentice-Hall. 455

Pu, C. (1988). Superdatabases for composition of heterogeneous databases. In *Proc. 4th Int. Conf. on Data Engineering*, pages 548–555. 147, 352

Pu, C. and Leff, A. (1991). Replica control in distributed systems: An asynchronous approach. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 377–386. 462

Pugh, W. (1989). Skip lists: A probabilistic alternative to balanced trees. In *Proc. Workshop on Algorithms and Data Structures*, pages 437–449. 622

Qiao, L., Agrawal, D., and Abbadi, A. E. (2003). Supporting sliding window queries for continuous data streams. In *Proc. 15th Int. Conf. on Scientific and Statistical Database Management*, pages 85–94. 737

Raghavan, S. and Garcia-Molina, H. (2001). Crawling the hidden web. In *Proc. 27th Int. Conf. on Very Large Data Bases*, pages 129–138. 657, 686

Raghavan, S. and Garcia-Molina, H. (2003). Representing web graphs. In *Proc. 19th Int. Conf. on Data Engineering*, pages 405–416. 658, 661, 662, 663

Rahal, A., Zhu, Q., and Larson, P.-Å. (2004). Evolutionary techniques for updating query cost models in a dynamic multidatabase environment. *VLDB J.*, 13(2):162–176. 307, 313, 331

Rahimi, S. (1987). Reference architecture for distributed database management systems. In *Proc. 3th Int. Conf. on Data Engineering*. Tutorial Notes. 40

Rahm, E. and Bernstein, P. A. (2001). A survey of approaches to automatic schema matching. *VLDB J.*, 10(4):334–350. 138, 139, 143, 146, 160

Rahm, E. and Do, H. H. (2000). Data cleaning: Problems and current approaches. *Q. Bull. IEEE TC on Data Eng.*, 23(4):3–13. 157

Rahm, E. and Marek, R. (1995). Dynamic multi-resource load balancing in parallel database systems. In *Proc. 21th Int. Conf. on Very Large Data Bases*, pages 395–406. 530, 548

Ramabhadran, S., Ratnasamy, S., Hellerstein, J. M., and Shenker, S. (2004). Brief announcement: prefix hash tree. In *Proc. ACM SIGACT-SIGOPS 23rd Symp. on the Principles of Distributed Computing*, page 368. 622, 643

Ramakrishnan, R. (2009). Data management in the cloud. In *Proc. 25th Int. Conf. on Data Engineering*, page 5. 753, 763

Ramakrishnan, R. and Gehrke, J. (2003). *Database Management Systems*. McGraw-Hill, 3 edition. 70, 189, 201

Ramamoorthy, C. V. and Wah, B. W. (1983). The isomorphism of simple file allocation. *IEEE Trans. Comput.*, C-23(3):221–231. 121

Ramamritham, K. and Pu, C. (1995). A formal characterization of epsilon serializability. *IEEE Trans. Knowl. and Data Eng.*, 7(6):997–1007. 401, 462

Raman, V., Deshpande, A., and Hellerstein, J. M. (2003). Using state modules for adaptive query processing. In *Proc. 19th Int. Conf. on Data Engineering*, pages 353–365. 331

Raman, V. and Hellerstein, J. M. (2001). Potter's wheel: An interactive data cleaning system. In *Proc. 27th Int. Conf. on Very Large Data Bases*, pages 381–390. 158

Ramanathan, P. and Shin, K. G. (1988). Checkpointing and rollback recovery in a distributed system using common time base. In *Proc. 7th Symp. on Reliable Distributed Systems*, pages 13–21. 456

Randell, B., Lee, P. A., and Treleaven, P. C. (1978). Reliability issues in computing system design. *ACM Comput. Surv.*, 10(2):123–165. 406, 455

Rao, P. and Moon, B. (2004). Prix: Indexing and querying XML using prüfer sequences. In *Proc. 20th Int. Conf. on Data Engineering*, pages 288–300. 701

Ratnasamy, S., Francis, P., Handley, M., and Karp, R. (2001a). A scalable content-addressable network. In *Proc. ACM Int. Conf. on Data Communication*, pages 161–172. 620, 646

Ratnasamy, S., Francis, P., Handley, M., Karp, R. M., and Shenker, S. (2001b). A scalable content-addressable network. In *Proc. ACM Int. Conf. on Data Communication*, pages 161–172. 618

Ray, I., Mancini, L. V., Jajodia, S., and Bertino, E. (2000). Asep: A secure and flexible commit protocol for mls distributed database systems. *IEEE Trans. Knowl. and Data Eng.*, 12(6):880–899. 187, 202

Reiss, F. and Hellerstein, J. (2005). Data triage: an adaptive architecture for load shedding in telegraphCQ. In *Proc. 21st Int. Conf. on Data Engineering*, pages 155–156. 740

Ribeiro-Neto, B. A. and Barbosa, R. A. (1998). Query performance for tightly coupled distributed digital libraries. In *Proc. 3rd ACM Int. Conf. on Digital Libraries*, pages 182–190. 668

Ritter, J. Why Gnutella can't scale, no, really (2001). Available from: `http://www.darkridge.com/~jpr5/doc/gnutella.html` [Last retrieved: December 2009]. 618

Rivera-Vega, P., Varadarajan, R., and Navathe, S. B. (1990). Scheduling data redistribution in distributed databases. In *Proc. Int. Conf. on Data Eng*, pages 166–173. 124

Rivest, R. L., Shamir, A., and Adelman, L. (1978). A method for obtaining digital signatures and public-key cryptosystems. *Commun. ACM*, 21(2):120–126. 180

Rjaibi, W. (2004). An introduction to multilevel secure relational database management systems. In *Proc. Conf. of the IBM Centre for Advanced Studies on Collaborative Research*, pages 232–241. 187, 202

Röhm, U., Böhm, K., and Schek, H.-J. (2000). Olap query routing and physical design in a database cluster. In *Advances in Database Technology, Proc. 7th Int. Conf. on Extending Database Technology*, pages 254–268. 535, 544, 548

Röhm, U., Böhm, K., and Schek, H.-J. (2001). Cache-aware query routing in a cluster of databases. In *Proc. 17th Int. Conf. on Data Engineering*, pages 641–650. 535

Röhm, U., Böhm, K., Schek, H.-J., and Schuldt, H. (2002a). Fas - a freshness-sensitive coordination cocoon for a cluster of olap components. In *Proc. 28th Int. Conf. on Very Large Data Bases*, pages 754–765. 493

Röhm, U., Böhm, K., Schek, H.-J., and Schuldt, H. (2002b). FAS - A freshness-sensitive coordination middleware for a cluster of olap components. In *Proc. 28th Int. Conf. on Very Large Data Bases*, pages 754–765. 462, 541

Roitman, H. and Gal, A. (2006). Ontobuilder: Fully automatic extraction and consolidation of ontologies from web sources using sequence semantics. In *EDBT Workshops*, volume 4254 of *LNCS*, pages 573–576. 152

Rosenkrantz, D. J. and Hunt, H. B. (1980). Processing conjunctive predicates and queries. In *Proc. 6th Int. Conf. on Very Data Bases*, pages 64–72. 224, 241

Rosenkrantz, D. J., Stearns, R. E., and Lewis, P. M. (1978). System level concurrency control for distributed database systems. *ACM Trans. Database Syst.*, 3(2):178–198. 390

Roth, J. P., Bouricius, W. G., Carter, E. C., and Schneider, P. R. (1967). Phase ii of an architectural study for a self-repairing computer. Report SAMSO-TR-67-106, U. S. Air Force Space and Missile Division, El Segundo, Calif. Cited in [Siewiorek and Swarz, 1982]. 410

Roth, M. and Schwartz, P. (1997). Don't scrap it, wrap it! a wrapper architecture for legacy data sources. In *Proc. 23th Int. Conf. on Very Large Data Bases*, pages 266–275. 327

Roth, M. T., Ozcan, F., and Haas, L. M. (1999). Cost models do matter: Providing cost information for diverse data sources in a federated system. In *Proc. 25th Int. Conf. on Very Large Data Bases*, pages 599–610. 307, 310, 331

Rothermel, K. and Mohan, C. (1989). Aries/nt: A recovery method based on write-ahead logging for nested transactions. In *Proc. 15th Int. Conf. on Very Large Data Bases*, pages 337–346. 401

Rothnie, J. B. and Goodman, N. (1977). A survey of research and development in distributed database management. In *Proc. 3rd Int. Conf. on Very Data Bases*, pages 48–62. 116

Rowstron, A. I. T. and Druschel, P. (2001). Pastry: Scalable, decentralized object location, and routing for large-scale peer-to-peer systems. In *Proc. IFIP/ACM Int. Conf. on Distributed Systems Platforms*, pages 329–350. 621

Ryvkina, E., Maskey, A., Adams, I., Sandler, B., Fuchs, C., Cherniack, M., and Zdonik, S. (2006). Revision processing in a stream processing engine: A high-level design. In *Proc. 22nd Int. Conf. on Data Engineering*, page 141. 725

Sacca, D. and Wiederhold, G. (1985). Database partitioning in a cluster of processors. *ACM Trans. Database Syst.*, 10(1):29–56. 99, 115, 125

Sacco, M. S. and Yao, S. B. (1982). Query optimization in distributed data base systems. In Yovits, M., editor, *Advances in Computers*, volume 21, pages 225–273. Academic Press. 39, 209, 211, 220

Saito, Y. and Shapiro, M. (2005). Optimistic replication. *ACM Comput. Surv.*, 37(1):42–81. 462, 466, 493

Salton, G. (1989). *Automatic Text Processing – The Transformation, Analysis, and Retrieval of Information by Computer*. Addison–Wesley. 667

Schlageter, G. and Dadam, P. (1980). Reconstruction of consistent global states in distributed databases. In Delobel, C. and Litwin, W., editors, *Distributed Data Bases*, pages 191–200. North-Holland. 456

Schlichting, R. D. and Schneider, F. B. (1983). Fail–stop processors: An approach to designing fault–tolerant computing systems. *ACM Trans. Comp. Syst.*, 1(3):222–238. 455

Schmidt, C. and Parashar, M. (2004). Enabling flexible queries with guarantees in p2p systems. *IEEE Internet Computing*, 8(3):19–26. 622

Schmidt, S., Berthold, H., and Legler, T. (2004). QStream: Deterministic querying of data streams. In *Proc. 30th Int. Conf. on Very Large Data Bases*, pages 1365–1368. 738

Schmidt, S., Legler, T., Schar, S., and Lehner, W. (2005). Robust real-time query processing with QStream. In *Proc. 31st Int. Conf. on Very Large Data Bases*, pages 1299–1301. 738

Schreiber, F. (1977). A framework for distributed database systems. In *Proc. Int. Computing Symposium*, pages 475–482. 39

Selinger, P. G. and Adiba, M. (1980). Access path selection in distributed data base management systems. In *Proc. First Int. Conf. on Data Bases*, pages 204–215. 250, 254, 277, 292, 293

Selinger, P. G., Astrahan, M. M., Chamberlin, D. D., Lorie, R. A., and Price, T. G. (1979). Access path selection in a relational database management system. In

Proc. ACM SIGMOD Int. Conf. on Management of Data, pages 23–34. 212, 253, 261, 292, 586

Serrano, D., Patiño-Martínez, M., Jiménez-Peris, R., and Kemme, B. (2007). Boosting database replication scalability through partial replication and 1-copy-snapshot-isolation. In *Proc. 13th IEEE Pacific Rim Int. Symp. on Dependable Computing*, pages 290–297. 491

Sevcik, K. C. (1983). Comparison of concurrency control methods using analytic models. In *Information Processing '83*, pages 847–858. 401

Severence, D. G. and Lohman, G. M. (1976). Differential files: Their application to the maintenance of large databases. *ACM Trans. Database Syst.*, 1(3):256–261. 419

Shafer, J. C., Agrawal, R., and Mehta, M. (1996). Sprint: A scalable parallel classifier for data mining. In *Proc. 22th Int. Conf. on Very Large Data Bases*, pages 544–555. 743

Shah, M. A., Hellerstein, J. M., Chandrasekaran, S., and Franklin, M. J. (2003). Flux: An adaptive partitioning operator for continuous query systems. In *Proc. 19th Int. Conf. on Data Engineering*, pages 25–36. 320, 321, 322, 331

Shapiro, L. (1986). oin processing in database systems with large main memories. *ACM Trans. Database Syst.*, 11(3):239–264. 587

Sharaf, M., Labrinidis, A., Chrysanthis, P., and Pruhs, K. (2005). Freshness-aware scheduling of continuous queries in the dynamic web. In *Proc. 8th Int. Workshop on the World Wide Web and Databases*, pages 73–78. 735

Sharp, J. (1987). *An Introduction to Distributed and Parallel Processing*. Blackwell Scientific Publications. 498

Shasha, D. and Wang, T.-L. (1991). Optimizing equijoin queries in distributed databases where relations are hash partitioned. *ACM Trans. Database Syst.*, 16(2):279–308. 292

Shatdal, A. and Naughton, J. F. (1993). Using shared virtual memory for parallel join processing. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 119–128. 534, 548

Shekita, E. J. and Carey, M. J. (1990). A performance evaluation of pointer-based joins. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 300–311. 590

Shekita, E. J., Young, H. C., and Tan, K. L. (1993). Multi-join optimization for symmetric multiprocessor. In *Proc. 19th Int. Conf. on Very Large Data Bases*, pages 479–492. 530, 548

Sheth, A. and Larson, J. (1990). Federated databases: Architectures and integration. *ACM Comput. Surv.*, 22(3):183–236. 40, 135, 160, 298

Sheth, A., Larson, J., Cornellio, A., and Navathe, S. B. (1988a). A tool for integrating conceptual schemas and user views. In *Proc. 4th Int. Conf. on Data Engineering*, pages 176–183. 147, 202

Sheth, A., Larson, J., and Watkins, E. (1988b). Tailor, a tool for updating views. In *Advances in Database Technology, Proc. 1st Int. Conf. on Extending Database Technology*, pages 190–213. Springer. 202

Sheth, A. P. and Kashyap, V. (1992). So far (schematically) yet so near (semantically). In *Proc. IFIP WG 2.6 Database Semantics Conf. on Interoperable Database Systems*, pages 283–312. 141

Shivakumar, N. and García-Molina, H. (1997). Wave-indices: indexing evolving databases. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 381–392. 738

Shrivastava, S. K., editor (1985). *Reliable Computer Systems*. Springer. 455, 768

Sidell, J., Aoki, P. M., Sah, A., Staelin, C., Stonebraker, M., and Yu, A. (1996). Data replication in mariposa. In *Proc. 12th Int. Conf. on Data Eng*, pages 485–494. 456, 493

Siegel, J., editor (1996). *CORBA Fundamentals and Programming*. John Wiley & Sons. 607

Siewiorek, D. P. and Swarz, R. S., editors (1982). *The Theory and Practice of Reliable System Design*. Digital Press. 407, 409, 455, 810

Silberschatz, A., Korth, H., and Sudarshan, S. (2002). *Database System Concepts*. McGraw-Hill, 4 edition. 70

Simon, E. and Valduriez, P. (1984). Design and implementation of an extendible integrity subsystem. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 9–17. 193, 202

Simon, E. and Valduriez, P. (1986). Integrity control in distributed database systems. In *Proc. 19th Hawaii Int. Conf. on System Sciences*, pages 622–632. 192, 202

Simon, E. and Valduriez, P. (1987). Design and analysis of a relational integrity subsystem. Technical Report DB-015-87, Microelectronics and Computer Corporation, Austin, Tex. 189, 192, 202

Singhal, M. (1989). Deadlock detection in distributed systems. *Comp.*, 22(11):37–48. 401

Sinha, M. K., Nanadikar, P. D., and Mehndiratta, S. L. (1985). Timestamp based certification schemes for transactions in distributed database systems. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 402–411. 385

Skarra, A. (1989). oncurrency control for cooperating transactions in an object-oriented database. In *Proc. ACM SIGPLAN Workshop on Object-Based Concurrent Programming*, pages 145–147. 401

Skarra, A., Zdonik, S., and Reiss, S. (1986). An object server for an object-oriented database system. In *Proc. of the 1st Int. Workshop on Object-Oriented Database Systems*, pages 196–204. 401

Skeen, D. (1981). Nonblocking commit protocols. In *ACM SIGMOD Int. Conf. on Management of Data*, pages 133–142. 440, 443, 447, 456

Skeen, D. (1982a). *Crash Recovery in a Distributed Database Management System*. Ph.D. thesis, Department of Electrical Engineering and Computer Science, University of California at Berkeley, Berkeley, Calif. 456

Skeen, D. (1982b). A quorum-based commit protocol. In *Proc. 6th Berkeley Workshop on Distributed Data Management and Computer Networks*, pages 69–80. 448, 450

Skeen, D. and Stonebraker, M. (1983). A formal model of crash recovery in a distributed system. *IEEE Trans. Softw. Eng.*, SE-9(3):219–228. 437, 443, 449, 456

Skeen, D. and Wright, D. (1984). Increasing availability in partitioned networks. In *Proc. 3rd ACM SIGACT–SIGMOD Symp. on Principles of Database Systems*, pages 290–299. 456, 493

Smith, J. M. and Chang, P. Y. (1975). Optimizing the performance of a relational algebra database interface. *Commun. ACM*, 18(10):568–579. 228, 241

Somani, A., Choy, D., and Kleewein, J. C. (2002). Bringing together content and data management systems: Challenges and opportunities. *IBM Systems J.*, 41(4):686–696. 159

Sousa, A., Oliveira, R., Moura, F., and Pedone, F. (2001). Partial replication in the database state machine. In *Proc. IEEE Int. Symp. Network Computing and Applications*, pages 298–309. 491, 548

Srivastava, U. and Widom, J. (2004). Memory-limited execution of windowed stream joins. In *Proc. 30th Int. Conf. on Very Large Data Bases*, pages 324–335. 740

Stallings, W. (2011). *Data and Computer Communications*. Prentice-Hall, 9 edition. 70

Stanoi, I., Agrawal, D., and El-Abbadi, A. (1998). Using broadcast primitives in replicated databases. In *Proc. 8th Int. Conf. on Distributed Computing Systems*, pages 148–155. 482

Stearns, R. E., II, P. M. L., and Rosenkrantz, D. J. (1976). Concurrency controls for database systems. In *Proc. 17th Symp. on Foundations of Computer Science*, pages 19–32. 350

Stöhr, T., Märtens, H., and Rahm, E. (2000). Multi-dimensional database allocation for parallel data warehouses. In *Proc. 26th Int. Conf. on Very Large Data Bases*, pages 273–284. 542

Stoica, I., Morris, R., Karger, D. R., Kaashoek, M. F., and Balakrishnan, H. (2001a). Chord: A scalable peer-to-peer lookup service for internet applications. In *Proc. ACM Int. Conf. on Data Communication*, pages 149–160. 618

Stoica, I., Morris, R., Liben-Nowell, D., Karger, D., Kaashoek, M., Dabek, F., and Balakrishnan, H. (2001b). Chord: A scalable peer-to-peer lookup protocol for internet applications. In *Proc. ACM Int. Conf. on Data Communication*, pages 149–160. 621

Stonebraker, M. (1975). Implementation of integrity constraints and views by query modification. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 65–78. 172, 173, 186, 191, 192, 201, 202

Stonebraker, M. (1981). Operating system support for database management. *Commun. ACM*, 24(7):412–418. 39, 415

Stonebraker, M. (1986). The case for shared nothing. *Q. Bull. IEEE TC on Data Eng.*, 9(1):4–9. 547

Stonebraker, M. (2010). SQL databases v. NoSQL databases. *Commun. ACM*, 53(4):10–11. 753

Stonebraker, M., Abadi, D. J., DeWitt, D. J., Madden, S., Paulson, E., Pavlo, A., and Rasin, A. (2010). MapReduce and parallel DBMSs: friends or foes? *Commun. ACM*, 53(1):64–71. 760, 763

Stonebraker, M. and Brown, P. (1999). *Object-Relational DBMSs*. Morgan Kaufmann, 2nd edition. 552, 607

Stonebraker, M., Kreps, P., Wong, W., and Held, G. (1976). The design and implementation of ingres. *ACM Trans. Database Syst.*, 1(3):198–222. 56, 258

Stonebraker, M. and Neuhold, E. (1977). A distributed database version of ingres. In *Proc. 2nd Berkeley Workshop on Distributed Data Management and Computer Networks*, pages 9–36. 474

Stonebraker, M., Rowe, L., Lindsay, B., Gray, J., Carey, M., Brodie, M., Bernstein, P., and Beech, D. (1990). Third-generation data base system manifesto. *ACM SIGMOD Rec.*, 19(3):31–44. 553

Straube, D. and Özsu, M. T. (1990a). Queries and query processing in object-oriented database systems. *ACM Trans. Information Syst.*, 8(4):387–430. 585

Straube, D. and Özsu, M. T. (1990b). Type consistency of queries in an object-oriented database. In *Proc. Joint ACM OOPSLA/ECOOP '90 Conference on Object-Oriented Programming: Systems, Languages and Applications*, pages 224–233. 585

Straube, D. D. and Özsu, M. T. (1995). Query optimization and execution plan generation in object-oriented database systems. *IEEE Trans. Knowl. and Data Eng.*, 7(2):210–227. 589

Strong, H. R. and Dolev, D. (1983). Byzantine agreement. In *Digest of Papers — COMPCON*, pages 77–81, San Francisco, Calif. 456

Stroustrup, B. (1986). *The C++ Programming Language*. Addison Wesley. 559

Sullivan, M. and Heybey, A. (1998). Tribeca: A system for managing large databases of network traffic. In *Proc. USENIX 1998 Annual Technical Conf.* 726, 730

Swami, A. (1989). Optimization of large join queries: combining heuristics and combinatorial techniques. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 367–376. 212, 249

Tandem (1987). Nonstop sql – a distributed high-performance, high-availability implementation of sql. In *Proc. Int. Workshop on High Performance Transaction Systems*, pages 60–104. 377, 548

Tandem (1988). A benchmark of nonstop sql on the debit credit transaction. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 337–341. 377

Tanenbaum, A. (1995). *Distributed Operating Systems*. Prentice-Hall. 180

Tanenbaum, A. S. (2003). *Computer Networks*. Prentice-Hall, 4th edition. 60, 70

Tanenbaum, A. S. and van Renesse, R. (1988). Voting with ghosts. In *Proc. 8th Int. Conf. on Distributed Computing Systems*, pages 456–461. 493

Tanenbaum, A. S. and van Steen, M. (2002). *Distributed Systems: Principles and Paradigms*. Prentice-Hall. 2

Tao, Y. (2010). *Mining Time-Changing Data Streams*. PhD thesis, University of Waterloo. 763

Tao, Y. and Özsu, M. T. (2009). Efficient decision tree construction for mining time-varying data streams. In *Proc. Conf. of the IBM Centre for Advanced Studies on Collaborative Research*. 743

Tao, Y., Yiu, M. L., Papadias, D., Hadjieleftheriou, M., and Mamoulis, N. (2005). RPJ: Producing fast join results on streams through rate-based optimization. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 371–382. 738

Tatarinov, I., Ives, Z. G., Madhavan, J., Halevy, A. Y., Suciu, D., Dalvi, N. N., Dong, X., Kadiyska, Y., Miklau, G., and Mork, P. (2003). The piazza peer data management project. *ACM SIGMOD Rec.*, 32(3):47–52. 625, 654

Tatbul, N., Cetintemel, U., Zdonik, S., Cherniack, M., and Stonebraker, M. (2003). Load shedding in a data stream manager. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 309–320. 739

Terry, D., Goldberg, D., Nichols, D., and Oki, B. (1992). Continuous queries over append-only databases. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 321–330. 6

Thakkar, S. S. and Sweiger, M. (1990). Performance of an oltp application on symmetry multiprocessor system. In *Proc. 17th Int. Symposium on Computer Architecture*, pages 228–238. 503

Thiran, P., Hainaut, J.-L., Houben, G.-J., and Benslimane, D. (2006). Wrapper-based evolution of legacy information systems. *ACM Trans. Softw. Eng. and Methodology*, 15(4):329–359. 329, 331

Thomas, R. H. (1979). A majority consensus approach to concurrency control for multiple copy databases. *ACM Trans. Database Syst.*, 4(2):180–209. 385, 450, 487

Thomasian, A. (1993). Two-phase locking and its thrashing behavior. *ACM Trans. Database Syst.*, 18(4):579–625. 401

Thomasian, A. (1996). *Database Concurrency Control: Methods, Performance, and Analysis*. Kluwer Academic Publishers. 358, 398, 399, 401

Thomasian, A. (1998). Distributed optimistic concurrency control methods for high performance transaction processing. *IEEE Trans. Knowl. and Data Eng.*, 10(1):173–189. 401

Thuraisingham, B. (2001). Secure distributed database systems. *Information Security Technical Report*, 6(2). 187, 202

Tian, F. and DeWitt, D. (2003a). Tuple routing strategies for distributed Eddies. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 333–344. 739

Tian, F. and DeWitt, D. J. (2003b). Tuple routing strategies for distributed eddies. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 333–344. 322, 326, 331

Tomasic, A., Amouroux, R., Bonnet, P., Kapitskaia, O., Naacke, H., and Raschid, L. (1997). The distributed information search component (DISCO) and the world-wide web – prototype demonstration. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 546–548. 319, 329

Tomasic, A., Raschid, L., and Valduriez, P. (1996). Scaling heterogeneous databases and the design of disco. In *Proc. 16th Int. Conf. on Distributed Computing Systems*, pages 449–457. 319, 331

Tomasic, A., Raschid, L., and Valduriez, P. (1998). Scaling access to distributed heterogeneous data sources with Disco. In *IEEE Trans. Knowl. and Data Eng.* in press. 319, 331

Traiger, I. L., Gray, J., Galtieri, C. A., and Lindsay, B. G. (1982). Transactions and recovery in distributed database systems. *ACM Trans. Database Syst.*, 7(3):323–342. 456

Triantafillou, P. and Pitoura, T. (2003). Towards a unifying framework for complex query processing over structured peer-to-peer data networks. In *Int. Workshop on Databases, Information Systems and Peer-to-Peer Computing*, pages 169–183. 641

Triantafillou, P. and Taylor, D. J. (1995). The location-based paradigm for replication: Achieving efficiency and availability in distributed systems. *IEEE Trans. Softw. Eng.*, 21(1):1–18. 493

Tsichritzis, D. and Klug, A. (1978). The ansi/x3/sparc dbms framework report of the study group on database management systems. *Inf. Syst.*, 1:173–191. 22

Tsuchiya, M., Mariani, M. P., and Brom, J. D. (1986). Distributed database management model and validation. *IEEE Trans. Softw. Eng.*, SE-12(4):511–520. 401

Tucker, P., Maier, D., Sheard, T., and Faragas, L. (2003). Exploiting punctuation semantics in continuous data streams. *IEEE Trans. Knowl. and Data Eng.*, 15(3):555–568. 725, 732

Ullman, J. (1997). Information integration using logical views. In *Proc. 6th Int. Conf. on Database Theory*, volume 1186 of *Lecture Notes in Computer Science*, pages 19–40. Springer. 303, 331

Ullman, J. D. (1982). *Principles of Database Systems*. Computer Science Press, 2nd edition. 224, 228, 231, 241, 272

Ullman, J. D. (1988). *Principles of Database and Knowledge Base Systems*, volume 1. Computer Science Press. 300, 301, 337

Ulusoy, Ö. (2007). Research issues in peer-to-peer data management. In *Proc. 22nd Int. Symp. on Computer and Information Science*, pages 1–8. 653

Urhan, T. and Franklin, M. J. (2000). XJoin: A reactively-scheduled pipelined join operator. *Q. Bull. IEEE TC on Data Eng.*, 23(2):27–33. 732

Urhan, T. and Franklin, M. J. (2001). Dynamic pipeline scheduling for improving interactive query performance. In *Proc. 27th Int. Conf. on Very Large Data Bases*, pages 501–510. 738

Urhan, T., Franklin, M. J., and Amsaleg, L. (1998a). Cost based query scrambling for initial delays. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 130–141. 322, 331

Urhan, T., Franklin, M. J., and Amsaleg, L. (1998b). Cost-based query scrambling for initial delays. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 130–141. 739

Valduriez, P. (1982). Semi-join algorithms for distributed database machines. In Schneider, J.-J., editor, *Distributed Data Bases*. North-Holland. pages 23–37. 270, 273, 291, 292

Valduriez, P. (1987). Join indices. *ACM Trans. Database Syst.*, 12(2):218–246. 587, 588, 589

Valduriez, P. (1993). Parallel database systems: Open problems and new issues. *Distrib. Parall. Databases*, 1:137–16. 497

Valduriez, P. and Boral, H. (1986). Evaluation of recursive queries using join indices. In *Proc. First Int. Conf. on Expert Database Systems*, pages 197–208. 219

Valduriez, P. and Gardarin, G. (1984). Join and semi-join algorithms for a multi processor database machine. *ACM Trans. Database Syst.*, 9(1):133–161. 291, 292, 513

Valduriez, P., Khoshafian, S., and Copeland, G. (1986). Implementation techniques of complex objects. In *Proc. 11th Int. Conf. on Very Large Data Bases*, pages 101–109. 579

Valduriez, P. and Pacitti, E. (2004). Data management in large-scale p2p systems. In *Proc. 6th Int. Conf. High Performance Comp. for Computational Sci.*, pages 104–118. 612, 653

Varadarajan, R., Rivera-Vega, P., and Navathe, S. B. (1989). Data redistribution scheduling in fully connected networks. In *Proc. 27th Annual Allerton Conf. on Communication, Control, and Computing*. 124

Velegrakis, Y., Miller, R. J., and Popa, L. (2004). Preserving mapping consistency under schema changes. *VLDB J.*, 13(3):274–293. 156, 157

Verhofstadt, J. S. (1978). Recovery techniques for database systems. *ACM Comput. Surv.*, 10(2):168–195. 39, 419, 456

Vermeer, M. (1997). *Semantic Interoperability for Legacy Databases*. Ph.D. thesis, Department of Computer Science, University of Twente, Enschede, Netherlands. 140

Vidal, M.-E., Raschid, L., and Gruser, J.-R. (1998). A meta-wrapper for scaling up to multiple autonomous distributed information sources. In *Proc. Int. Conf. on Cooperative Information Systems*, pages 148–157. 314

Viglas, S. and Naughton, J. (2002). Rate-based query optimization for streaming information sources. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 37–48. 738, 739

Viglas, S., Naughton, J., and Burger, J. (2003). Maximizing the output rate of multi-join queries over streaming information sources. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 285–296. 732, 739

Vossough, E. and Getta, J. R. (2002). Processing of continuous queries over unlimited data streams. In *Proc. 13th Int. Conf. Database and Expert Systems Appl.*, pages 799–809. 733

Voulgaris, S., Jelasity, M., and van Steen, M. (2003). A robust and scalable peer-to-peer gossiping protocol. In *Agents and Peer-to-Peer Computing, Second Int. Workshop, (AP2PC)*, pages 47–58. 618

Vu, Q. H., Lupu, M., and Ooi, B. C. (2009). *Peer-to-Peer Computing: Principles and Applications*. Springer. 653

Wah, B. W. and Lien, Y. N. (1985). Design of distributed databases on local computer systems. *IEEE Trans. Softw. Eng.*, SE-11(7):609–619. 214, 215

Walsh, N., editor. The DocBook schema (2006). Available from: `http://www.oasis-open.org/docbook/specs/wd-docbook-docbook-5.0b3.html` [Last retrieved: December 2009]. 690

Walton, C., Dale, A., and Jenevin, R. (1991). A taxonomy and performance model of data skew effects in parallel joins. In *Proc. 17th Int. Conf. on Very Large Data Bases*, pages 537–548. 527, 548

Wang, H., Fan, W., Yu, P., and Han, J. (2003a). Mining concept-drifting data streams using ensemble classifiers. In *Proc. 9th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, pages 226–235. 743

Wang, H. and Meng, X. (2005). On the sequencing of tree structures for XML indexing. In *Proc. 21st Int. Conf. on Data Engineering*, pages 372–383. 701

Wang, H., Park, S., Fan, W., and Yu, P. S. (2003b). ViST: A dynamic index method for querying XML data by tree structures. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 110–121. 701

Wang, H., Zaniolo, C., and Luo, R. (2003c). Atlas: A small but complete SQL extension for data mining and data streams. In *Proc. 29th Int. Conf. on Very Large Data Bases*, pages 1113–1116. 732

Wang, S., Rundensteiner, E., Ganguly, S., and Bhatnagar, S. (2006). State-slice: New paradigm of multi-query optimization of window-based stream queries. In *Proc. 32nd Int. Conf. on Very Large Data Bases*. 740

Wang, W., Li, J., Zhang, D., and Guo, L. (2004). Processing sliding window join aggregate in continuous queries over data streams. In *Proc. 8th East European Conf. Advances in Databases and Information Systems*, pages 348–363. 733

Wang, Y. and Rowe, L. (1991). Cache consistency and concurrency control in a client/server dbms architecture. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 367–376. 573

Weihl, W. (1988). Commutativity-based concurrency control for abstract data types. *IEEE Trans. Comput.*, C-37(12):1488–1505. 594, 595, 604

Weihl, W. (1989). Local atomicity properties: Modular concurrency control for abstract data types. *ACM Trans. Prog. Lang. and Syst.*, 11(2):249–28. 594, 595

Weikum, G. (1986). Pros and cons of operating system transactions for data base systems. In *Proc. AFIPS Fall Joint Computer Conf.*, pages 1219–1225. 397

Weikum, G. (1991). Principles and realization strategies of multilevel transaction management. *ACM Trans. Database Syst.*, 16(1):132–180. 397, 398

Weikum, G. and Hasse, C. (1993). Multi-level transaction management for complex objects: Implementation, performance, parallelism. *VLDB J.*, 2(4):407–454. 397, 604, 605

Weikum, G. and Schek, H. J. (1984). Architectural issues of transaction management in layered systems. In *Proc. 10th Int. Conf. on Very Large Data Bases*, pages 454–465. 397

Weikum, G. and Vossen, G. (2001). *Transactional Information Systems: Theory, Algorithms, and the Practice of Concurrency Control*. Morgan Kaufmann. 358

White, S. and DeWitt, D. (1992). Quickstore: A high performance mapped object store. In *Proc. 18th Int. Conf. on Very Large Data Bases*, pages 419–431. 576

Wiederhold, G. (1982). *Database Design*. McGraw-Hill, 2nd edition. 83

Wiederhold, G. (1992). Mediators in the architecture of future information systems. *Comp.*, 25(3):38–49. 37, 331

Wiesmann, M., Schiper, A., Pedone, F., Kemme, B., and Alonso, G. (2000). Database replication techniques: A three parameter classification. In *Proc. 19th Symp. on Reliable Distributed Systems*, pages 206–215. 493

Wilkinson, K. and Neimat, M. (1990). Maintaining consistency of client-cached data. In *Proc. 16th Int. Conf. on Very Large Data Bases*, pages 122–133. 572

Williams, R., Daniels, D., Haas, L., Lapis, G., Lindsay, B., Ng, P., Obermarck, R., Selinger, P., Walker, A., Wilms, P., and Yost, R. (1982). R*: An overview of the architecture. In *Proc. 2nd Int. Conf. on Databases*, pages 1–28. 175, 214, 215

Wilms, P. F. and Lindsay, B. G. (1981). A database authorization mechanism supporting individual and group authorization. Research Report RJ 3137, IBM Almaden Research Laboratory, San Jose, Calif. 186, 187, 201

Wilschut, A. and Apers, P. (1991). Dataflow query execution in a parallel main-memory environment. In *Proc. 1st Int. Conf. on Parallel and Distributed Information Systems*, pages 68–77. 322, 325, 641, 732

Wilshut, A. N. and Apers, P. (1992). Parallelism in a main-memory system: The performance of prisma/db. In *Proc. 22th Int. Conf. on Very Large Data Bases*, pages 23–27. 526

Wilshut, A. N., Flokstra, J., and Apers, P. (1995). Parallel evaluation of multi-join queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 115–126. 529, 534

Wilson, B. and Navathe, S. B. (1986). An analytical framework for the redesign of distributed databases. In *Proc. 6th Advanced Database Symposium*, pages 77–83. 124

Wolf, J. L., Dias, D., Yu, S., and Turek, J. (1993). Algorithms for parallelizing relational database joins in the presence of data skew. Research Report RC19236 (83710), IBM Watson Research Center, Yorktown Heights, NY. 529, 548

Wolfson, O. (1987). The overhead of locking (and commit) protocols in distributed databases. *ACM Trans. Database Syst.*, 12(3):453–471. 455, 456, 493

Wong, E. (1977). Retrieving dispersed data from sdd-1. In *Proc. 2nd Berkeley Workshop on Distributed Data Management and Computer Networks*, pages 217–235. 281, 293

Wong, E. and Youssefi, K. (1976). Decomposition: A strategy for query processing. *ACM Trans. Database Syst.*, 1(3):223–241. 258, 275, 292

Wright, D. D. (1983). Managing distributed databases in partitioned networks. Technical Report TR83-572, Department of Computer Science, Cornell University, Ithaca, N.Y. 456, 493

Wu, E., Diao, Y., and Rizvi, S. (2006). High-performance complex event processing over streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 407–418. 725

Wu, K.-L., Chen, S.-K., and Yu, P. (2004). Interval query indexing for efficient stream processing. In *Proc. 13th ACM Int. Conf. on Information and Knowledge Management*, pages 88–97. 741

Wu, K.-L., Yu, P. S., and Pu, C. (1997). Divergence control algorithms for epsilon serializability. *IEEE Trans. Knowl. and Data Eng.*, 9(2):262–274. 401, 462

Wu, S., Yu, G., Yu, Y., Ou, Z., Yang, X., and Gu, Y. (2005). A deadline-sensitive approach for real-time processing of sliding windows. In *Proc. 6th Int. Conf. on Web-Age Information Management:*, pages 566–577. 740

Fernández, M., Malhotra, A., Marsh, J., Nagy, M., and Walsh, N., editors. XQuery 1.0 and XPath 2.0 data model (XDM) (2007). Available from: http://www.w3.org/TR/2007/REC-xpath-datamodel-20070123 [Last retrieved: February 2010]. 712

XHTML. XHTML 1.0 The extensible HyperText markup language (2nd edition) (2002). Available from: http://www.w3.org/TR/xhtml1/ [Last retrieved: December 2009]. 690

Xie, J., Yang, J., and Chen, Y. (2005). On joining and caching stochastic streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 359–370. 740

Xu, J., Lin, X., and Zhou, X. (2004). Space efficient quantile summary for constrained sliding windows on a data stream. In *Proc. 5th Int. Conf. on Web-Age Information Management:*, pages 34–44. 737

Yan, L. L. (1997). Towards efficient and scalable mediation: The aurora approach. In *Proc. IBM CASCON Conference*, pages 15–29. 134

Yan, L.-L., Miller, R. J., Haas, L. M., and Fagin, R. (2001). Data-driven understanding and refinement of schema mappings. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 485–496. 152

Yan, L.-L. and Özsu, M. T. (1999). Conflict tolerant queries in aurora. In *Proc. Int. Conf. on Cooperative Information Systems*, pages 279–290. 158

Yan, L. L., Özsu, M. T., and Liu, L. (1997). Accessing heterogeneous data through homogenization and integration mediators. In *Proc. Int. Conf. on Cooperative Information Systems*, pages 130–139. 134

Yang, B. and Garcia-Molina, H. (2002). Improving search in peer-to-peer networks. In *Proc. 22nd Int. Conf. on Distributed Computing Systems*, pages 5–14. 617

Yang, X., Lee, M.-L., and Ling, T. W. (2003). Resolving structural conflicts in the integration of XML schemas: A semantic approach. In *Proc. 22nd Int. Conf. on Conceptual Modeling*, pages 520–533. 134

Yao, S. B., Navathe, S. B., and Weldon, J.-L. (1982a). *An Integrated Approach to Database Design*, pages 1–30. Lecture Notes in Computer Science 132. Springer. 73

Yao, S. B., Waddle, V., and Housel, B. (1982b). View modeling and integration using the functional data model. *IEEE Trans. Softw. Eng.*, SE-8(6):544–554. 149

Yeung, C. and Hung, S. (1995). A new deadlock detection algorithm for distributed real-time database systems. In *Proc. 14th Symp. on Reliable Distributed Systems*, pages 146–153. 401

Yong, V., Naughton, J., and Yu, J. (1994). Storage reclamation and reorganization in client-server persistent object stores. In *Proc. 10th Int. Conf. on Data Engineering*, pages 120–133. 581

Yormark, B. (1977). The ansi/sparc/dbms architecture. In Jardine, D. A., editor, *ANSI/SPARC DBMS Model*, pages 1–21. North-Holland. 22

Yoshida, M., Mizumachi, K., Wakino, A., Oyake, I., and Matsushita, Y. (1985). Time and cost evaluation schemes of multiple copies of data in distributed database systems. *IEEE Trans. Softw. Eng.*, SE-11(9):954–958. 124

Yu, C. and Meng, W. (1998). *Principles of Query Processing for Advanced Database Applictions*. Morgan Kaufmann. 331

Yu, C. T. and Chang, C. C. (1984). Distributed query processing. *ACM Comput. Surv.*, 16(4):399–433. 220

Yu, P. S., Cornell, D., Dias, D. M., and Thomasian, A. (1989). Performance comparison of the io shipping and database call shipping schemes in multi-system partitioned database systems. *Perf. Eval.*, 10:15–33. 401

Zaniolo, C. (1983). The database language gem. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 207–218. 587

Zdonik, S. and Maier, D., editors (1990). *Readings in Object-Oriented Database Systems*. Morgan Kaufmann. 607

Zezula, P., Amato, G., Debole, F., and Rabitti, F. (2003). Tree signatures for XML querying and navigation. In *Database and XML Technologies, 1st Int. XML Database Symp.*, pages 149–163. 701

Zhang, C., Naughton, J. F., DeWitt, D. J., Luo, Q., and Lohman, G. M. (2001). On supporting containment queries in relational database management systems. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 425–436. 699, 700

Zhang, J. and Honeyman, P. (2008). A replicated file system for grid computing. *Concurrency and Computation: Practice and Experience*, 20(9):1113–1130. 750

Zhang, N. (2006). *Query Processing and Optimization in Native XML Databases*. PhD thesis, University of Waterloo. 719

Zhang, N., Agarwal, N., Chandrasekar, S., Idicula, S., Medi, V., Petride, S., and Sthanikam, B. (2009a). Binary XML storage and query processing in oracle 11g. *PVLDB*, 2(2):1354–1365. 703

Zhang, N., Kacholia, V., and Özsu, M. T. (2004). A succinct physical storage scheme for efficient evaluation of path queries in XML. In *Proc. 20th Int. Conf. on Data Engineering*, pages 54–65. 699

Zhang, N. and Özsu, M. T. (2010). XML native storage and query processing. In Li, C. and Ling, T.-W., editors, *Advanced Applications and Structures in XML Processing: Label Streams, Semantics Utilization and Data Query Technologies*. IGI Global. 699

Zhang, N., Özsu, M. T., Aboulnaga, A., and Ilyas, I. F. (2006a). XSEED: accurate and fast cardinality estimation for XPath queries. In *Proc. 22nd Int. Conf. on Data Engineering*, page 61. 702

Zhang, N., Özsu, M. T., Ilyas, I. F., and Aboulnaga, A. (2006b). Fix: Feature-based indexing technique for XML documents. In *Proc. 32nd Int. Conf. on Very Large Data Bases*, pages 259–270. 701

Zhang, R., Koudas, N., Ooi, B. C., and Srivastava, D. (2005). Multiple aggregations over data streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 299–310. 740

Zhang, Y. (2010). *XRPC: Efficient Distributed Query Processing on Heterogeneous XQuery Engines*. PhD thesis, Universiteit van Amsterdam. 719

Zhang, Y. and Boncz, P. A. (2007). Xrpc: Interoperable and efficient distributed XQuery. In *Proc. 33rd Int. Conf. on Very Large Data Bases*, pages 99–110. 712

Zhang, Y., Tang, N., and Boncz, P. A. (2009b). Efficient distribution of full-fledged XQuery. In *Proc. 25th Int. Conf. on Data Engineering*, pages 565–576. 710, 712

Zhao, B., Huang, L., Stribling, J., Rhea, S., Joseph, A. D., and Kubiatowicz, J. (2004). Tapestry: A resilient global-scale overlay for service deployment. *IEEE J. Selected Areas in Comm.*, 22(1):41–53. 620, 646

Zhu, Q. (1995). *Estimating Local Cost Parameters for Global Query Optimization in a Multidatabase System*. Ph.D. thesis, Department of Computer Science, University of Waterloo, Waterloo, Canada. 313

Zhu, Q. and Larson, P.-Å. (1994). A query sampling method of estimating local cost parameters in a multidatabase system. In *Proc. 10th Int. Conf. on Data Engineering*, pages 144–153. 307, 308, 331

Zhu, Q. and Larson, P. A. (1996a). Developing regression cost models for multidatabase systems. In *Proc. 4th Int. Conf. on Parallel and Distributed Information Systems*, pages 220–231. 307, 309, 331

Zhu, Q. and Larson, P. A. (1996b). Global query processing and optimization in the cords multidatabase system. In *Proc. Int. Conf. on Parallel and Distributed Computing Systems*, pages 640–647. 308

Zhu, Q. and Larson, P. A. (1998). Solving local cost estimation problem for global query optimization in multidatabase systems. *Distrib. Parall. Databases*, 6(4):373–420. 307, 308, 331

Zhu, Q. and Larson, P.-Å. (2000). Classifying local queries for global query optimization in multidatabase systems. *Int. J. Cooperative Information Syst.*, 9(3):315–355. 309

Zhu, Q., Motheramgari, S., and Sun, Y. (2003). Cost estimation for queries experiencing multiple contention states in dynamic multidatabase environments. *Knowledge and Information Systems*, 5(1):26–49. 307, 314, 331

Zhu, Q., Sun, Y., and Motheramgari, S. (2000). Developing cost models with qualitative variables for dynamic multidatabase environments. In *Proc. 16th Int. Conf. on Data Engineering*, pages 413–424. 307, 313, 331

Zhu, S. and Ravishankar, C. (2004). A scalable approach to approximating aggregate queries over intermittent streams. In *Proc. 16th Int. Conf. on Scientific and Statistical Database Management*, pages 85–94. 727

Zhu, Y., Rundensteiner, E., and Heineman, G. (2004). Dynamic plan migration for continuous queries over data streams. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, pages 431–442. 739

Zhu, Y. and Shasha, D. (2003). Efficient elastic burst detection in data streams. In *Proc. 9th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, pages 336–345. 727

Ziane, M., Zaït, M., and Borla-Salamet, P. (1993). Parallel query processing with zigzag trees. *VLDB J.*, 2(3):277–301. 523, 548

Zloof, M. M. (1977). Query-by-example: A data base language. *IBM Systems J.*, 16(4):324–343. 57

Zobel, D. D. (1983). The deadlock problem: A classifying bibliography. *Operating Systems Rev.*, 17(2):6–15. 401

Index